

Domain Logic Patterns (Patterns of Enterprise Application Architecture - Chapter 9)

Mike Helmick Large Scale Software Engineering - Spring 2014 University of Cincinnati

Patterns (Ch 9)

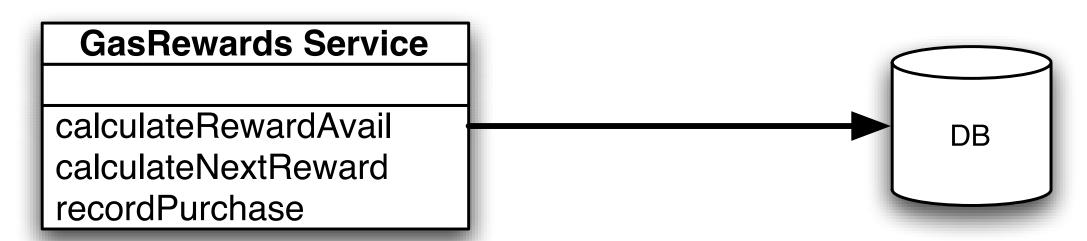


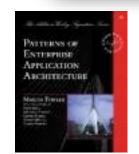
Patterns (Ch 9)

- Transaction Script (Page 110)
- Domain Model (Page 116)
- Table Module (Page 125)
- Service Layer (Page 133)



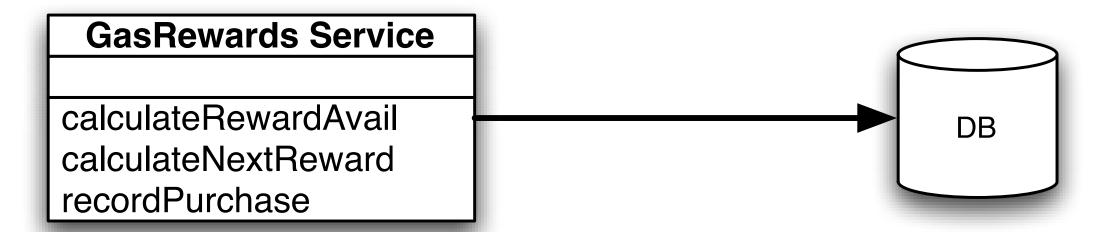








 "Organizes business logic by procedures where each procedure handles a single request from the presentation"









- All logic is organized in these large single procedures
- Will usually call database directly or use a thin wrapper



How it works



How it works

- When discovering your software requirements think about the coarse steps of the applications
- Your transaction is all in one procedure so we don't really have to worry about what other transactions are doing



Where it goes

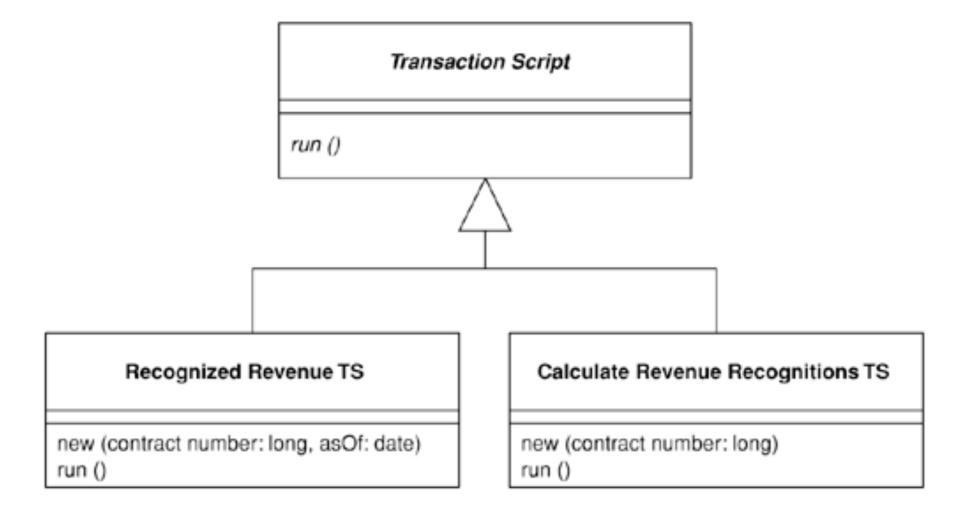


Where it goes

- Depends on your organizations
 - Server Page
 - CGI Script
 - Remote session object



Reuse

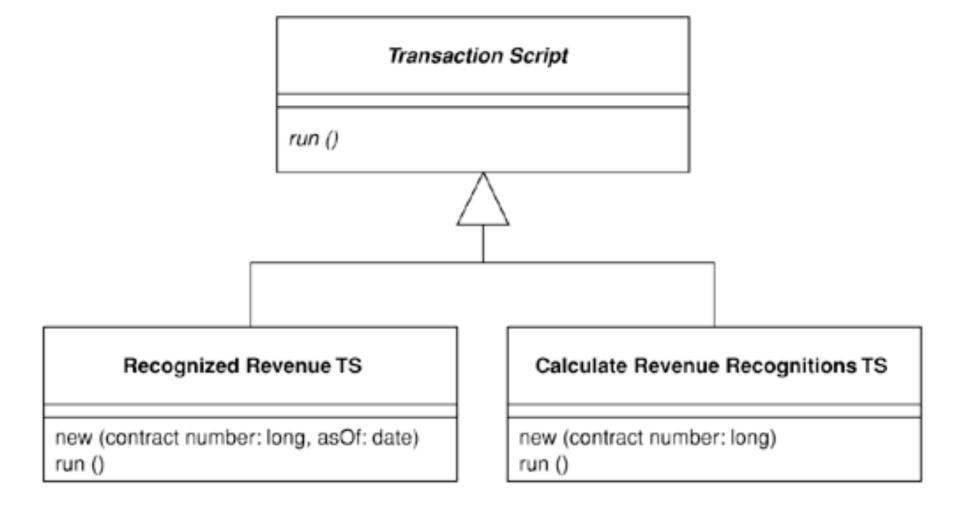






Reuse

- We can still organize for reuse
 - putting common functionality in a base class (transaction control / security)







Organization



Organization

- Several transaction script methods in the same class
 - Usually related business logic
- One class for each transaction script
 - Can take advantage of the "command pattern"



When to use



When to use

- When simplicity is important
- Small applications simple logic
- Have very low overhead possibly use in extremely performance critical applications



Drawbacks



Drawbacks

- Sort of becomes unwieldy as your application grows (think of having several hundred transaction script objects)
- A lot of code gets duplicated
- Difficult to refactor



Sample Problem





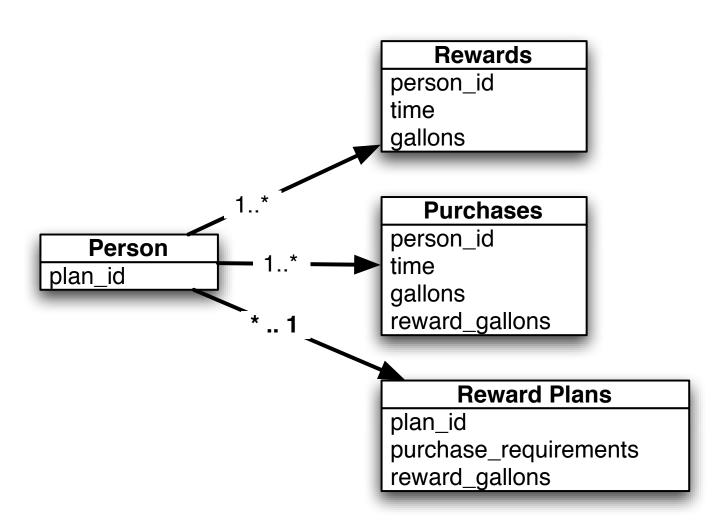
Sample Problem

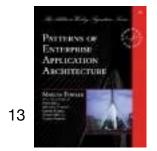
- Gas Rewards
 - Maybe we have different plans
 - Free gallon for every 10 purchased
 - Free gallon for every 15 purchased
 - Need to record purchases, calculate reward gallons, calculate what is needed for reward





Rules

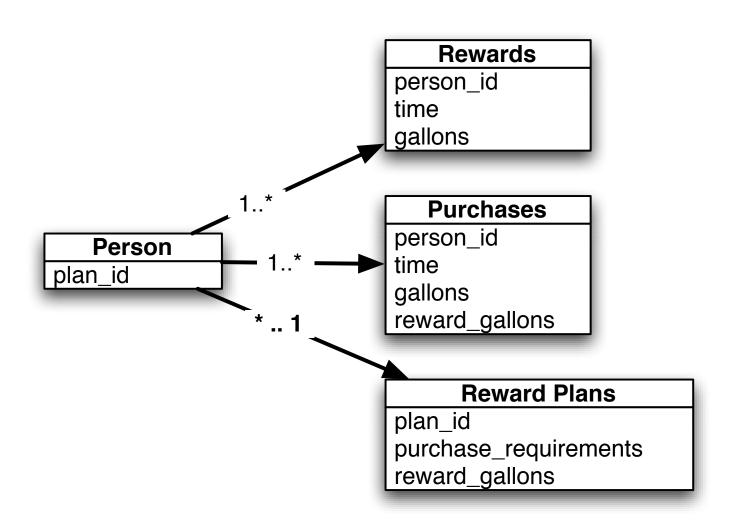


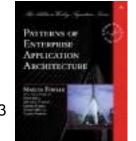




Rules

• Depending on the product purchased, revenue is recognized on different settings



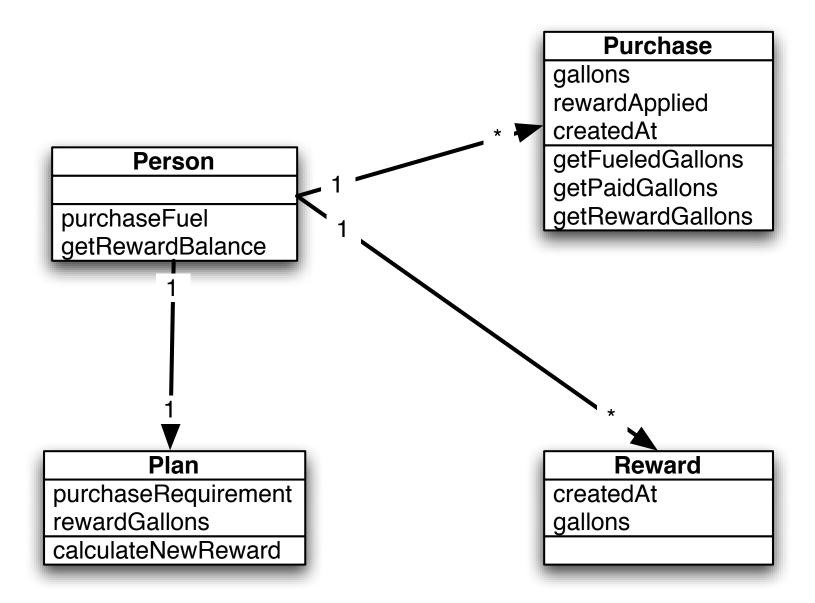




Example



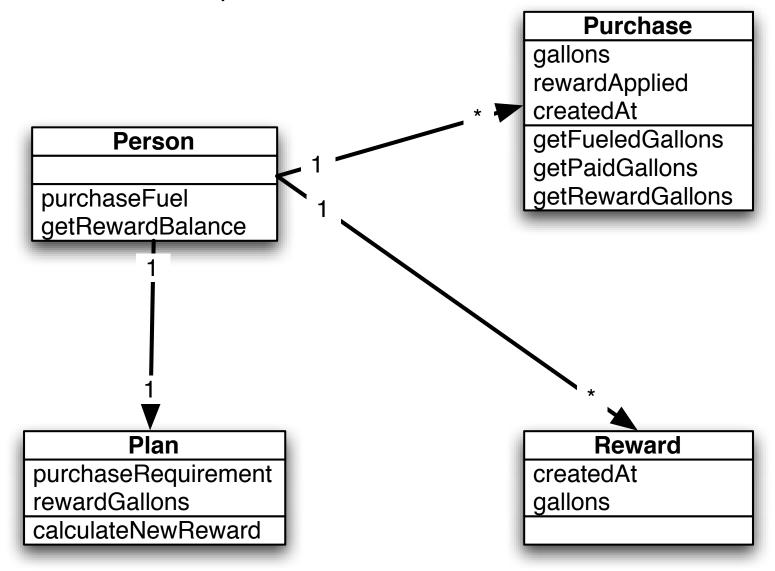








• "An object model of the domain that incorporates both behavior and data."









- Interconnected object where each object represents an individual entity within the system
- Very object-oriented
 - fits in well with



How it works



How it works

- You have an entire layer that is an object-oriented representation of the entire business logic (entities and actions)
- looks similar to database design
 - but different
 - Database only represents relationships
 - OO model does more





- Simple Domain Model
 - Very closely maps to the database
 - Easy to map to the database
- Rich Domain Model
 - Uses OO style for things inheritance, etc...
 - Difficult to map to the database



Mapping



Mapping

- Simple domain model
 - can use Active Record (advantages to this)
- Rich domain model
 - needs to use Data Mapper



Testing



Testing

- Minimize coupling between the Domain Model and the other layers
 - Specifically the ideal situation allows the Domain Model to function



Memory



Memory

- We can very rarely hold the entire object model in memory
- So
 - · we pull back the parts that we need from the database for the current calculation
 - we will even ignore related data until we need it (Lazy Load)



Bloat



Bloat

- Sometimes with Domain Model developers put too much one time use logic in the objects
- Is something is really specialized we can maybe move it to a service layer
- Have to balance between bloat and duplication



Java



Java

- As noted on page 118-119
- Using POJOs (Plain Old Java Objects) has many advantages
 - Easy to develop
 - No dependencies



When to Use



When to Use

- Depends on your complexity
- Domain Model scales very well (in terms of both complexity and number of entities)
- Learning how to design domain models takes lots of practice / experience



Considerations



Considerations

- Author prefers Data Mapper when using the Domain Model pattern
 - I prefer the Active Record pattern
- A Service Layer is an excellent addition since it gives your program an API
 - This is preferable to exposing your domain model & possibly allowing for violations of business logic



Example

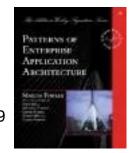








• "A single instance that handles the business logic for all rows in a database table or view."







- Organizes domain logic with 1 class per table
- Contains methods to act on the data for that table





- Works by packaging data and behavior (OO)
 - but more closely models the relational database
- What's missing
 - Identity the Table Module object has no identity





- Deals with tabular data (Record Set) which looks like the database table
- Could be an object that you instantiate or just a collection of static methods



Queries



Queries

- Can be built into the Table Module
 - or in a separate Table Data Gateway
- When you use a Table Data Gateway
 - adds complications
 - allows you to create a record set from scratch and test the Table Module



When to use



When to use

- When your data is tabular when it makes sense to operate on a whole table at the same time
- Doesn't represent relationships well
- Can't represent complex logic well
- Primarily used in .NET/COM



Example



Example

• See the C# example on pages 129 - 132







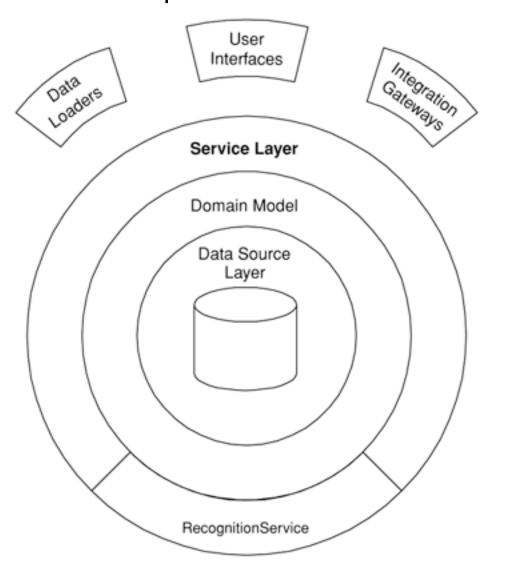


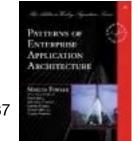
• "Defines an application's boundary with a layer of services that establishes a set of available operations and coordinates the application's response in each operation."



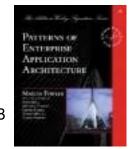


• "Defines an application's boundary with a layer of services that establishes a set of available operations and coordinates the application's response in each operation."



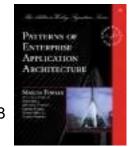








- Enterprise applications typically have multiple interfaces
 - GUI, Web Service, Batch Jobs, etc...







• There are many various ways to build a service interface



Kinds of Business Logic



Kinds of Business Logic

- There are two central kinds of logic
 - domain logic having to do with the specific problem at hand
 - application logic application responsibilities integration, notifications





- Domain Facade
 - The Service Layer manipulates the object-oriented Domain Model
 - Presents a well defined boundary between the layers of your application





- Operation script
 - Set of "thicker" classes that implement application logic
 - Still delegates domain logic to another layer



Remoting



Remoting

- Your service layer should be coarse grained
 - this makes it suitable for acting as a remote service layer
- Remote invocation is comparatively expensive



Remote



Remote

- If you provide a good separation of layers...
 - you can start coding locally and then distribute your service layer later



Identifying Services



Identifying Services

- The services you provide are typically defined by the requirements of the service consumers
 - i.e. if there are two comparable services we want the one thats easier to work with



Identifying



Identifying

- Most operations are CRUD type operations
- They still have lots of work that goes around them (validations, etc...)



Java



Java

- Can be implemented using POJOs or Stateless Session Beans
 - Ease of testing vs. transaction control
- Author recommends a hybrid approach
 - Stateless Session Beans delegating to POJOs
 - We can now control transactions without EJBs





Data Source Patterns (Patterns of Enterprise Application Architecture - Chapter 10)

Mike Helmick Large Scale Software Engineering - Spring 2014 University of Cincinnati

Patterns (Ch 10)

- Table Data Gateway
- Row Data Gateway
- Active Record
- Data Mapper





• "An object that acts as a Gateway to a database table. One instance handles all the rows in the table."

Person Gateway

find (id): RecordSet

findWithLastName(String) : RecordSet

update (id, lastname, firstname, numberOfDependents)

insert (lastname, firstname, numberOfDependents)

delete (id)







- Holds all of the SQL for the table
 - isolated from the application





- A very simple interface
 - Finder methods
 - Insert new record
 - Update existing
 - Delete record
- Usually stateless (could be static methods)





- All data is returned as multiple values (using a Record Set)
 - This requires you to know the underlying makeup of the data (raw database layout)
- You could modify this to return a Map based structure
 - Requires a pass through the mapping layer





- The Record Set still binds you to the specific data interface
- You can modify the Table Data Gateway to work with your Domain Model
 - This can cause problems with relationship mapping though
- Usually one class per database table





- First decide if a gateway should be used at all
- If there is a very simple database interface
 - easy to understand
 - quick to implement





- Works well with the Table Module data model
- Works well with Transaction Script
- Can also encapsulate stored procedures



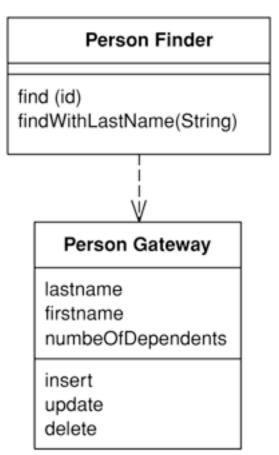
Example



Row Data Gateway



Row Data Gateway

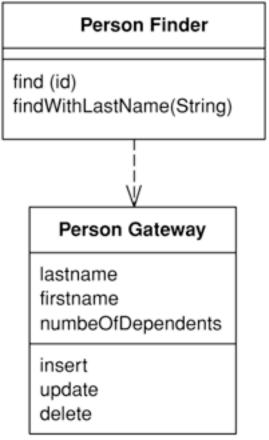






Row Data Gateway

• "An object that acts as a Gateway to a single record in a data source. There is one instance per row."









- An object that looks exactly like one row in the database (appropriate fields, etc)
- Does type conversion to/from database
- Only contains database access logic
 - & does not contain domain logic



Finders



Finders

- Can go in 2 locations
 - static method on the row gateway object
 - static method on "finder" class
- Putting them on a different class allows you to manage inheritance / polymorphism better



vs Active Record



vs Active Record

- The difference is
 - Active Record includes domain logic
 - Row Data Gateway does not include domain logic



Caution



Caution

- Row Data Gateway records are not created uniformly
 - There is no way to enforce that row '2' isn't loaded (and possibly manipulated) more than once
- They can interfere with each other



Writing



Writing

- Can take a long time to write
- Can be uniformly generated pretty easily
 - Code generation is something that you should always be thinking of as a possibility





- Again should you use a gateway at all?
- Works well with transaction script
- Doesn't work very well with Domain Model
 - Active Record does more with less code
 - Data Mapper works better if more complex





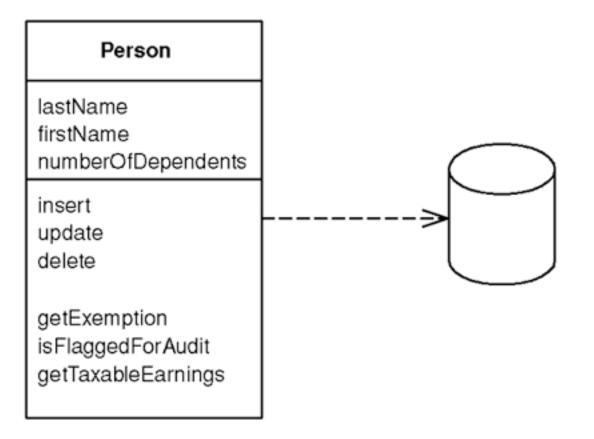
- Row Data Gateway provide good insulation from the database
 - You can even present different field names
 - You can aggregate some fields, etc...



Example



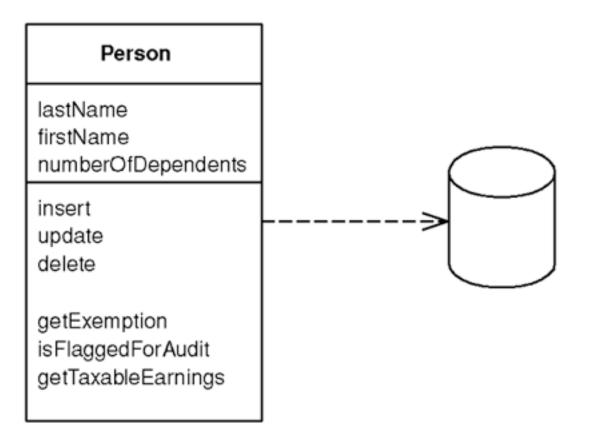


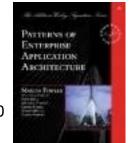






• "An object that wraps a row in a database table or view; encapsulates the database access, and adds domain logic on that data."









- A very obvious approach
- Every object is "activated" and knows how to create, find, update, and delete themselves





- Active Record is a Domain Model with classes that look very much like the database tables
 - Including field names, types, and relationships





- Each object
 - Has all of it's appropriate domain logic
 - Logic for loading/saving to the database



Fields



Fields

- Best if fields in the class match exactly to the fields in the database
 - This allows for automation







- "Construct an instance of the Active Record from a SQL result set row"
- "Construct a new instance for later insertion into the table"
- "Static finder methods to wrap commonly used SQL queries and return Active Record objects"
- "Update the database and insert into it the data in the Active Record"
- Get and set the fields
- Implement some pieces of business logic







- 'Get' methods
 - Can do SQL type on the fly
 - Can load dependent data (related Active Records either on creation or using Lazy Load)





- When domain logic isn't overly complex
 - most of the actions are CRUD type
 - this is usually the case





- Active Record is a simple
 - it is very easy to use
- But the objects must correspond directly to the database
 - Data Mapper works better when the logic / relationships are overly complex



Real Implementations



Real Implementations

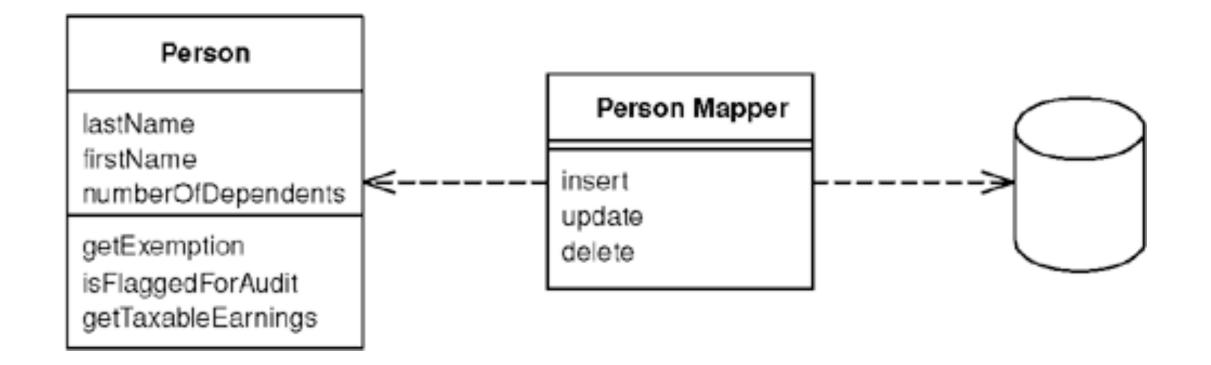
- This is what Ruby on Rails uses
 - And one of the primary reasons that development with Rails is so quick / easy
- Other languages now use this pattern in frameworks
 - Java ActiveJDBC https://github.com/javalite/activejdbc
 - Objective-C https://github.com/supermarin/ObjectiveRecord

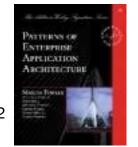


Example



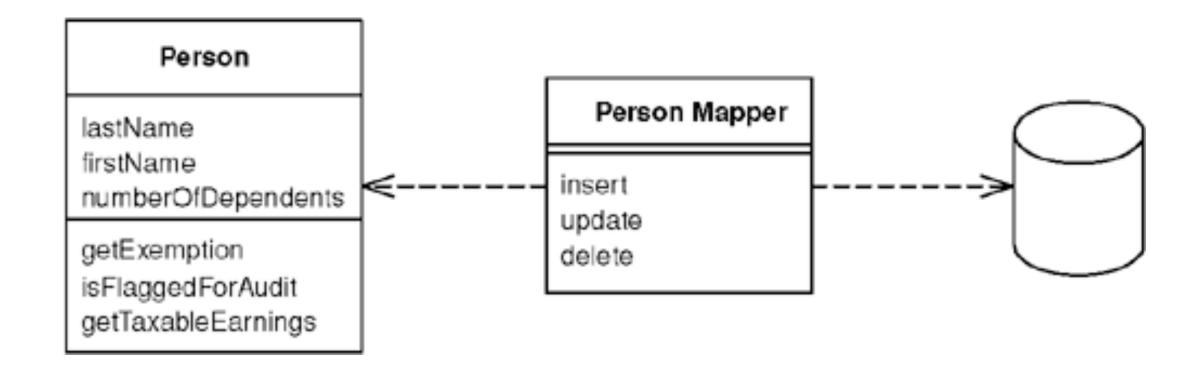








• "A layer of Mappers that moves data between objects and a database while keeping them independent of each other and the mapper itself."









- We've talked about things in OO technology that don't translate well to relational databases
- Coupled with complex business logic
 - We can use a Data Mapper to better organize things



Isolation



Isolation

- The in memory objects don't necessarily know about the database
- But -
 - changes in one will cause changes in the other





- The key is that the layers are separated
- Works well with caching mechanisms such as the Identity Map
- The Data Mapper layer can be swapped out for testing
 - or re-configured to point to a test database





- The Data Mapper needs to be intelligent enough to determine which fields of changed
- And to organize items into a transactional framework
 - Unit of Work is a candidate for this





- Lazy Load can be utilized
 - at a cost of making your Domain Model somewhat aware of the mapping layer





- Can be coded 2 ways
 - One mapper for each class when you hardcode them
 - One data mapper for everything when using Metadata Mapping



Finders



Finders

- Typical scenario: Service Layer finds a few objects and then kicks off domain logic
- The Domain Objects themselves usually need to find other objects
 - Can be handled with Lazy Load
 - Or an intermediary between the Domain Model and the Data Mapper





- The programming language created a problem
 - For some fields, we want them to be read-only to the user
 - But they must be writable by the Data Mapper





- A rich constructor can be used
 - A constructor which initializes the minimal set of fields required for a valid object
 - Created no invalid objects
 - Prevents immutable fields from being changed



Metadata Mapping



Metadata Mapping

- The simplest way to do mapping is to hard code everything
- Metadata mapping will store the database to domain model mapping in
 - a class, a file, the database itself
 - Can then either generate the mapper
 - or be used by the mapper





- When you want your object model and your database model to be isolated from each other
- Can potentially evolve them separately from each other
 - but that is not always the case





- An extra layer of software means that you have an extra layer to test
- If your business logic is simple, you can probably use active record
 - Then use Data Mapper for more complex business logic



Example





Object-Relational Behavioral Patterns (Patterns of Enterprise Application Architecture - Chapter 11)

Mike Helmick Large Scale Software Engineering University of Cincinnati Spring 2014

Patterns (Ch 11)



Patterns (Ch 11)

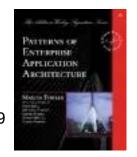
- Unit of Work
- Identity Map
- Lazy Load





Unit of Work

registerNew(object)
registerDirty (object)
registerClean(object)
registerDeleted(object)
commit()

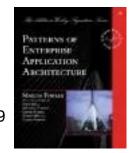




• "Maintains a list of objects affected by a business transaction and coordinates the writing out of changes and the resolution of concurrency problems"

Unit of Work

registerNew(object)
registerDirty (object)
registerClean(object)
registerDeleted(object)
commit()







Unit of Work

- Changes made to domain objects are useless if they can't be written back to the database
- We can make multiple SQL calls (each time something changes)
- or We can keep track of the changes





- Unit of Work is an object that keep track of
 - new, updated, and deleted objects
- The object handles inconsistent reads
- Tracks changed objects





- At commit time
 - The Unit of Work object decides what to do
 - Checks for concurrency violations
 - Pessimistic or Optimistic Offline Lock
- Application programmers never call methods that change the database



Registering



Registering

- An object can be registered with the Unit of Work
 - Manually by the developer
 - Automatically when the change
 - This obviously requires extra wiring
 - You could possibly do some interesting things with the Observer pattern



Caller client database a unit of work Registration new (ID) a customer select set credit limit register dirty (a customer) commit save update

Caller Registration

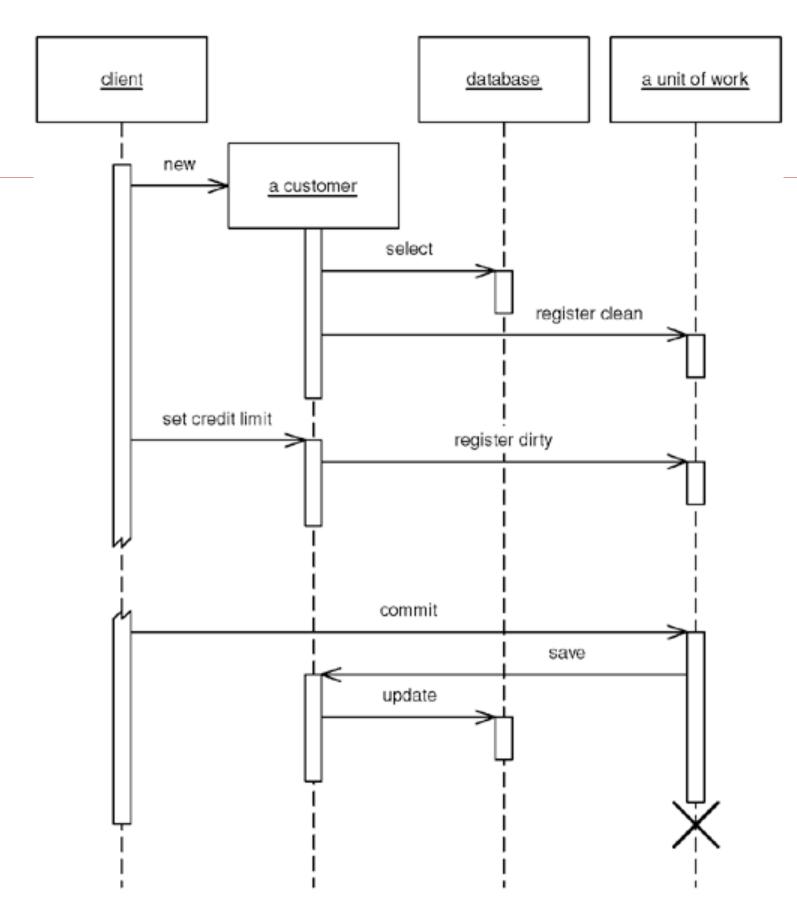


Caller Registration

- You have to remember to notify the Unit of Work
- Can lead to confusion
- Framework is easier to develop



Object Registration





Object Registration



Object Registration

- Place a registration call in the mutator methods of each domain object
- The Unit of Work needs to be passed to the object or in an easy to find place
 - Thread local storage is a candidate solution
 - Session object access can work as well



Aspects



Aspects

- The problem of inserting the calls to 'setDirty' is a candidate for Aspect Oriented Programming
- This can enforce that the method is called after every call to a 'set' method



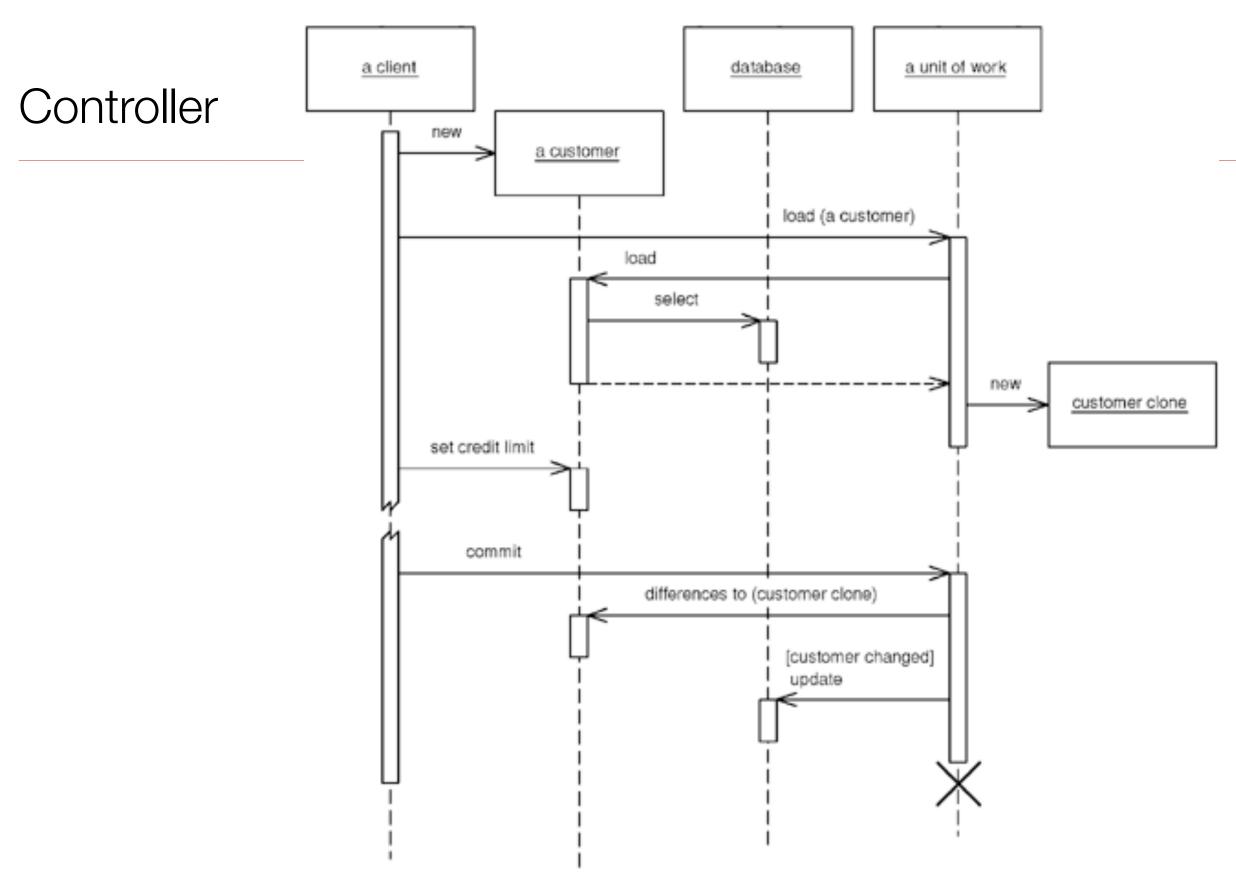
Controller



Controller

- Another option is the unit of work controller
- The Unit of Work handles all reads
 - The object is created twice
 - 1 copy goes to the application
 - 1 copy stays with the unit of work
 - Objects are then compared at commit time







Referential Integrity



Referential Integrity

- Unit of Work can solve this problem
 - Although your DB software might be able to solve it too
- You can define an absolute ordering to the tables & the objects can be saved in that order
 - Use metadata
 - A defined order can help eliminate deadlock



Batch Update



Batch Update

- A Unit of Work can also perform batch updates
- This is where SQL commands are grouped into a single database call
 - can be more efficient than separate calls for each statement



When to Use It



When to Use It

- When you want to combine database calls
- Usually the best way handle your updates from an efficiency



A Practical Matter



A Practical Matter

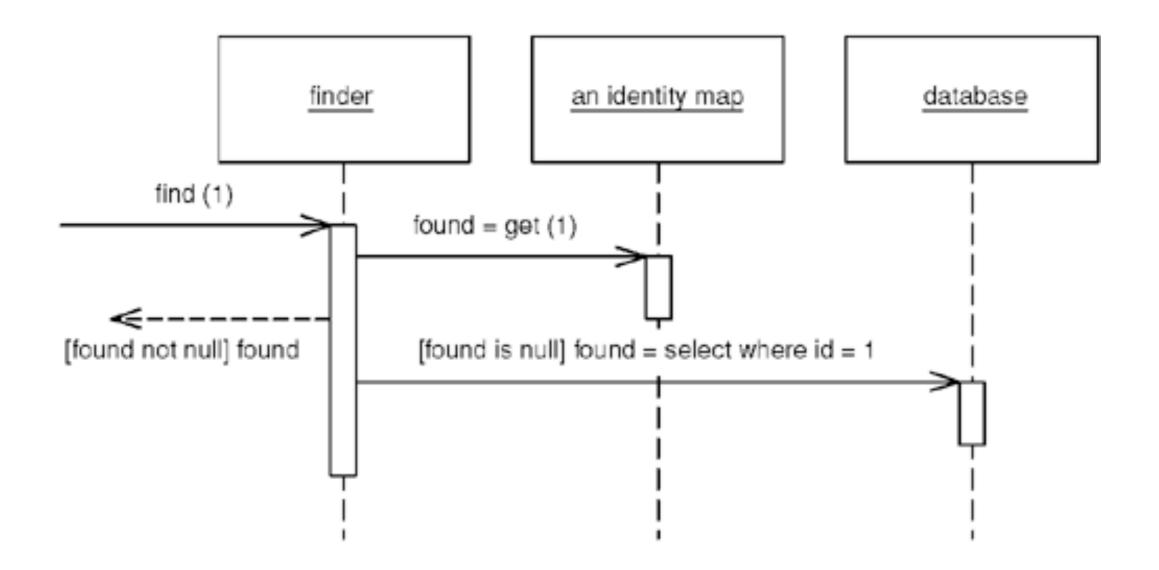
- This always sounds nice in theory to me
- I've never seen it used in a practical system
- Doesn't mesh well with certain data source layers

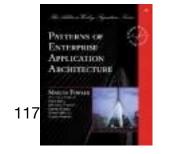


Example



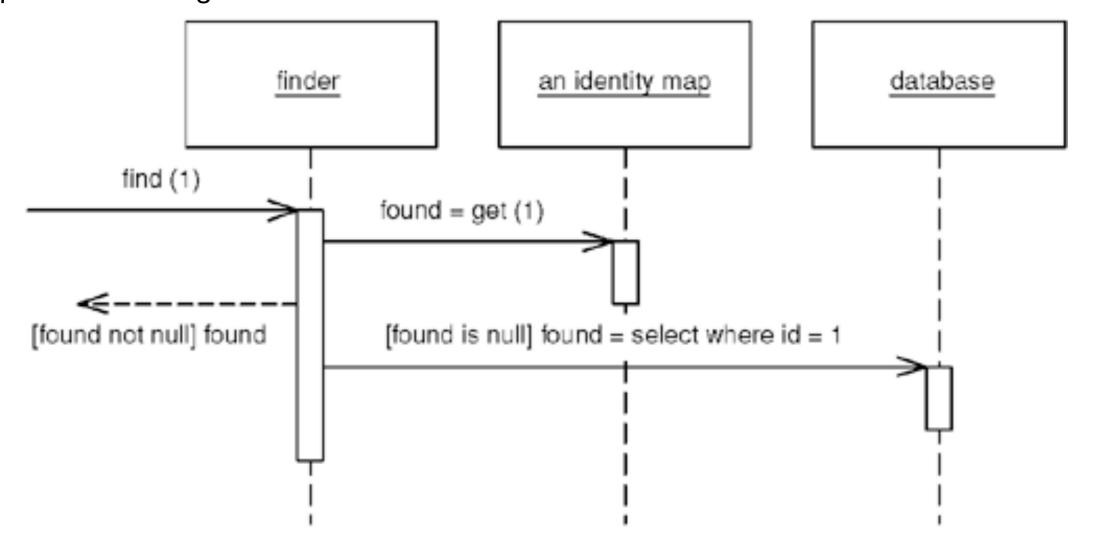








• "Ensures that each object gets loaded only once by keeping every loaded object in a map. Looks up objects using the map when referring to them."









- · Will make sure that you never load the same object from the database into memory twice
 - Can help performance
 - Cuts down on confusion (more important)





- Some kind of map structure to be the authoritative in-memory data source
- Can have one map per table or centralized maps, or whatever works for your schema



Keys



Keys

- You must choose a uniform key so that the identity map will function correctly
- Often a uniform ID column will work well



Explicit or Generic



Explicit or Generic

- You have to decide whether to use an explicit data structure or a generic data structure
 - Generic need to cast
 - Explicit no casing
 - This is much easier to do through Java Generics i.e. specialize your data structure



How Many



How Many

- One map per class
- One map for everything
- My thought
 - use a map of maps
 - or if you're using data mapper
 - Maybe you can put one in each mapper
 - or one in each active record





- The Identity Map has to be easy to access from multiple locations
- Identity Maps should be isolated on a per-session (or even per-request) basis
 - We can get creative though may some data is read-only and can be shared across all sessions





- From a practical standpoint
 - One Identity Map per transaction /request is probably the easiest to deal with
 - You can be sure to flush the identity map when appropriate



When to Use



When to Use

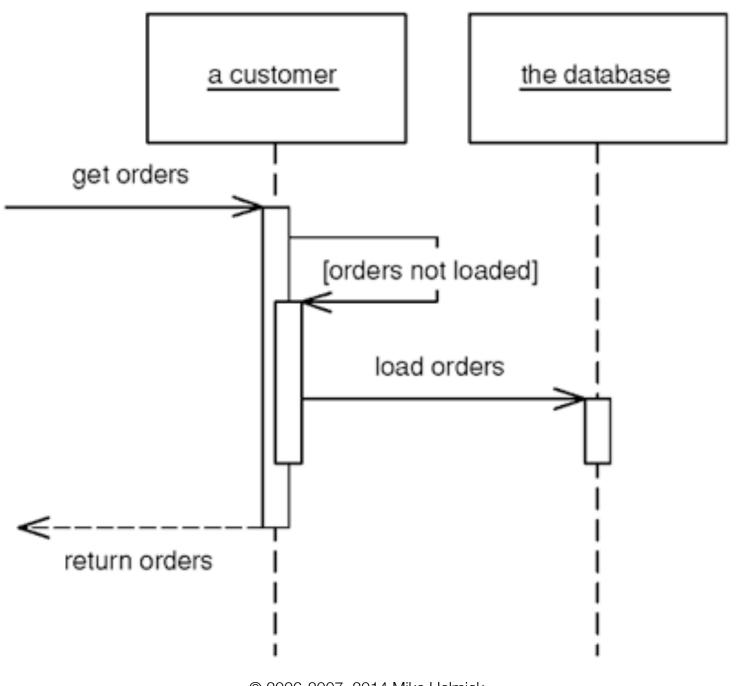
- When you want to be certain that you don't double load information from the database
- Might occur when the same record is read in multiple places during a transaction



Example



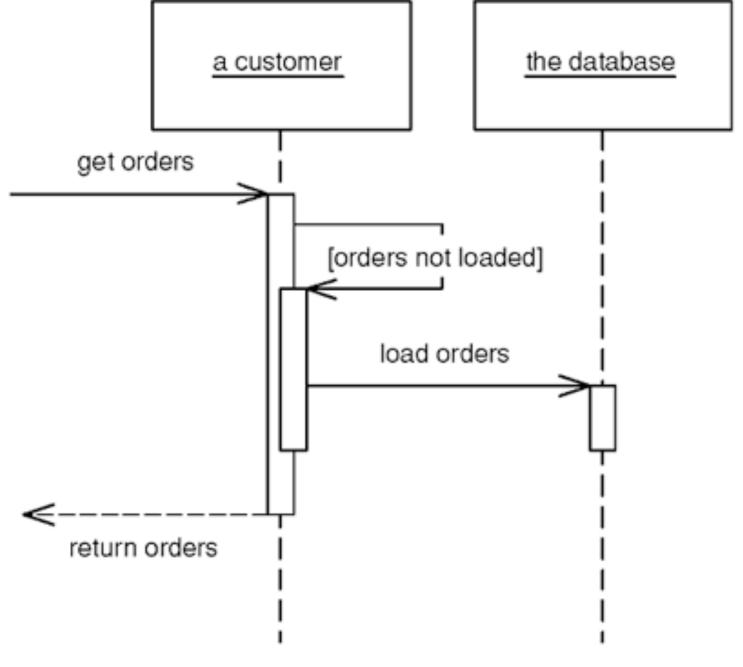


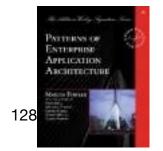






• "An object that doesn't contain all of the dat you need but knows how to get it."



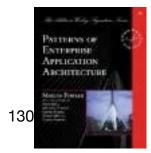




- The most obvious thing to do is to load and object and all related data into memory at the same time
 - A customer and all of their transactions
 - Over time this can lead to problems as the customer builds up more and more transactions
 - or sometimes we just want the address of the customer and don't need the transactions



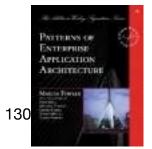
How it Works





How it Works

- 4 was to implement (from the book)
 - lazy initialization
 - virtual proxy
 - value holder
 - ghost







- Easiest to implement
- Each getter method that requires this must
 - first see if the field / collection / etc has been initialized
 - if not initialize it
 - return the value





- Usually null will work for detailing fields that haven't been loaded
 - This doesn't work if null is a valid value





- Can cause a dependency between the objects and the database
- Works best with Active Record, Table Data Gateway, and Row Data Gateway
- Can be made to work with Data Mapper
 - but works better with virtual proxy





- GOF Book Page 207
- An object that has the same interface as the object, but doesn't actually hold the data
- Both will implement the same interface





- Good
 - Looks identical to the real object
 - Easy to program
- Bad
 - There is an identity problem 2 in memory objects representing the same object



Value Holder



Value Holder

- An object that actually hold the real value
- But it only loads the value on the first access



Ghost



Ghost

- The object is created containing just the ID field
- When any other field is loaded
 - The rest of the object is loaded



Collections



Collections

- For performance reasons
 - · You will usually want to load an entire dependent collection of objects at the same time
 - Rather than loading the individual members as they are accessed



When to Use It



When to Use It

- Consider using lazy load if data would already require a second call to read the data
- Might as well delay the call in the event that the data ends up not being needed





Object-Relational Structural Patterns (Patterns of Enterprise Application Architecture - Chapter 12)

Mike Helmick Large Scale Software Engineering - Spring 2014 University of Cincinnati

Patterns (Ch 12)



Patterns (Ch 12)

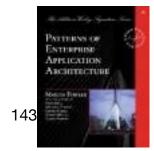
- Identity Field
- Foreign Key Mapping
- Association Table Mapping
- Dependent Mapping
- Embedded Value
- Serialized LOB
- Single Table Inheritance
- · Class Table Inheritance
- Concrete Table Inheritance
- Inheritance Mappers





Person

id: long

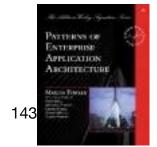




• "Saves a database ID field in an object to maintain identity between an in-memory object and a database row."

Person

id: long







- Database table require a primary key on every record to uniquely identify the records
- In memory objects do not have this requirement
- Identity Field bridges this gap



How It Works



How It Works

- The overall idea is very simple
 - We designate a key and use it to find rows by identity
 - Several complications come up



Key



Key

- Choosing the appropriate key is a difficult task
- Should we use naturally occurring keys or generated keys
 - Natural: A Student's SSN
 - Generated: A numeric ordering or hashed value



Natural Keys



Natural Keys

- Usually sound like a good idea
- Just doesn't work out that way all of the time
 - Mistyped SSN
 - Address (gets reused after a period of time)
 - Phone Number (same)



Key Width



Key Width

- A simple key is a single field that uniquely identified a record
- A compound key is when 2 or more fields are required to uniquely identify each record
 - Compound keys are more difficult to deal with
 - Can be avoided with generated keys
 - But they can be necessary with many-to-many relationships



Inheritance



Inheritance

- When using Concrete Table Inheritance or Class Table Inheritance
 - You must make sure the key is unique across several tables (the object hierarchy)





- In its simplest form it is a field that matches the type and name of the key on the database
 - private long id;



New Keys



New Keys

- Needed on all inserts / object creations
 - Auto-generate (From database)
 - Use GUID
 - Generate your own (possibly from database)



Auto-Generate



Auto-Generate

- Sounds easy
 - but it is database dependent
 - it is hard to get the generated value back right away



GUID



GUID

- Globally Unique IDentifier
 - A unique identifier that is guaranteed to be unique across all machines and times
 - Generates big keys (large numerical values) which can lead to performance and storage problems



Roll your own



Roll your own

- Use the SQL max function to find the highest
 - can severely impact performance
 - can lead to contention & race conditions
- Better approach
 - Use a key table



Key Table



Key Table

- A table with 2 fields, table name, last ID
 - Before each insert, read the last ID, add 1, and update
 - Best to have this in a separate transaction so as not to lock the key table
- Best choice for portability and efficiency



When to Use It



When to Use It

- When there is a direct mapping between domain model objects and database rows
 - Domain Model or Row Data Gateway
- Can assists with use of an Identity Map



Example



Example

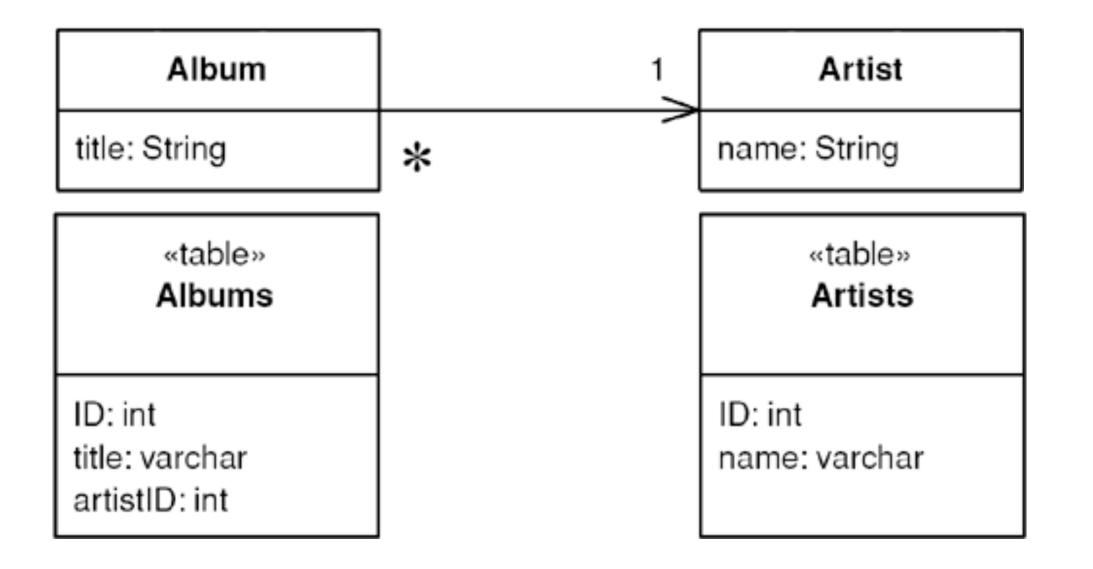
See pages 224 - 234 for an in-depth example using compound keys.



Foreign Key Mapping



Foreign Key Mapping

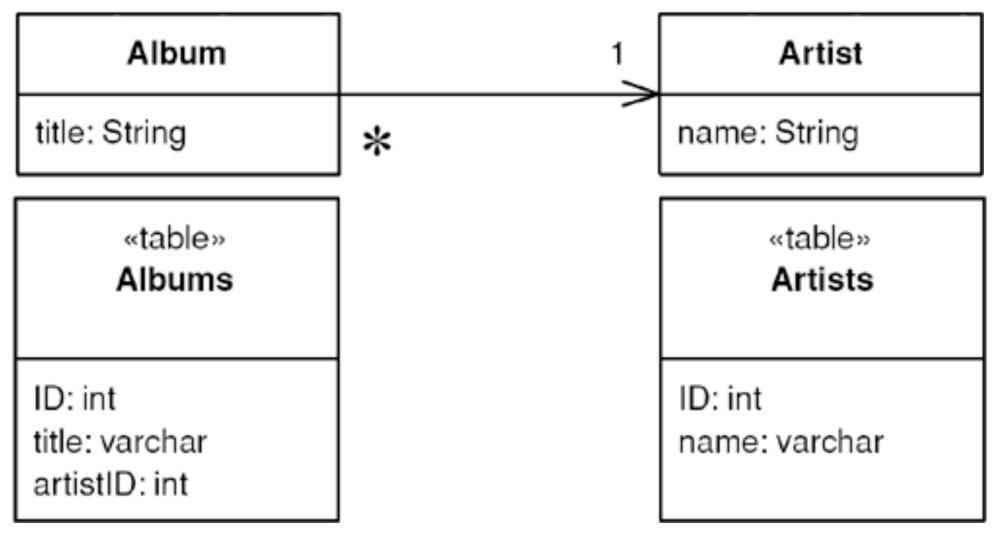






Foreign Key Mapping

• "Maps an association between objects to a foreign key reference between tables."







Foreign Key Mapping



Foreign Key Mapping

- In memory objects refer to each other seamlessly
 - Object references
- When we save/load to/from a database
 - Object references must be mapped to foreign key identities



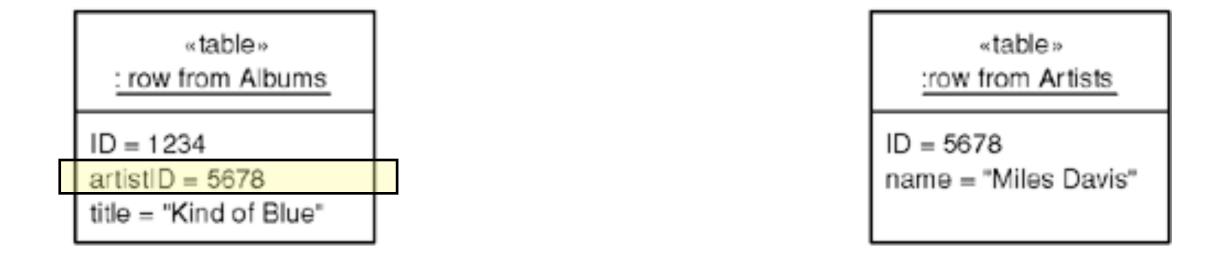


- Each record has an identity field of its own
- We must use the identity fields to link to the owning record / owned records
 - Usually the owning record



1 to 1

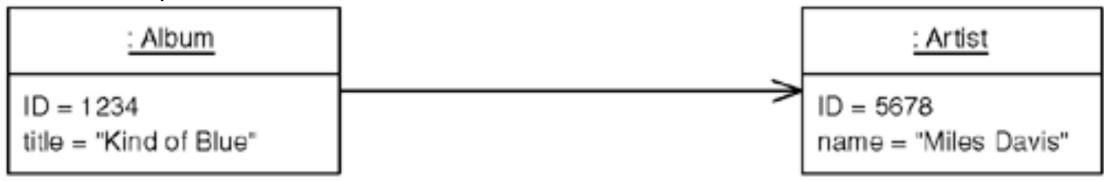






1 to 1

• With the 1 to 1 case the parent record can have the ID field of the child record







Collection

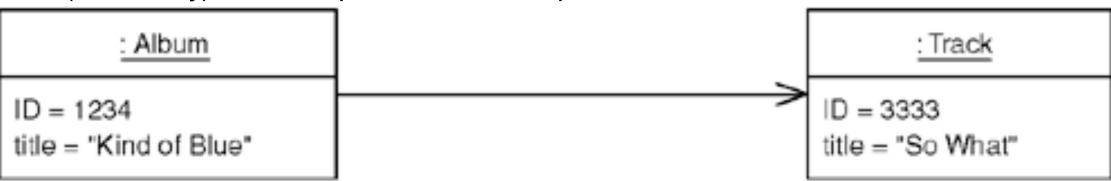






Collection

• With a collection (1 to many) we must put the ownership on the owned record









- Dependent items can be inserted / updated / removed
 - commits can get complicated





- "delete & insert"
 - delete all dependent record on the database
 - insert all the dependent records in memory
- Easy to code
- Not so great for performance reasons





- back pointers
 - give dependent records a reference back to the owning object
 - this makes it easier to set the ID field on the owned objects





- "Diff"
 - It is possible to compare the current state of the database to the list in memory
 - or copy when it is first read





- My recommendation
 - Use back-pointers
 - Save all the objects as if they exist on their own



Ownership



Ownership

- If we allow dependent records to be "orphaned"
 - Then code needs to be built to to handle knowing if a record is deleted (Because its back-pointer is nulled out) or just doesn't have an owner



Cycles



Cycles

- Links between objects are great
- When loading objects
 - be careful not to follow cycles in order to load the objects
 - Lazy load usually takes care of this since fields aren't loaded until accessed



When to Use It



When to Use It

- Can be used for almost all associations
 - many-to-many is the one exception



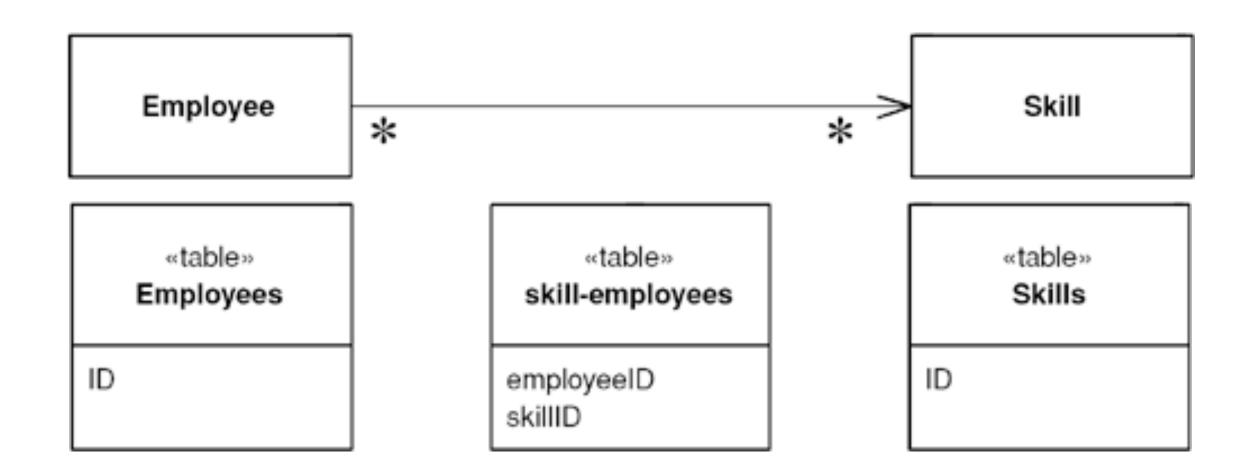
Example



Association Table Mapping



Association Table Mapping

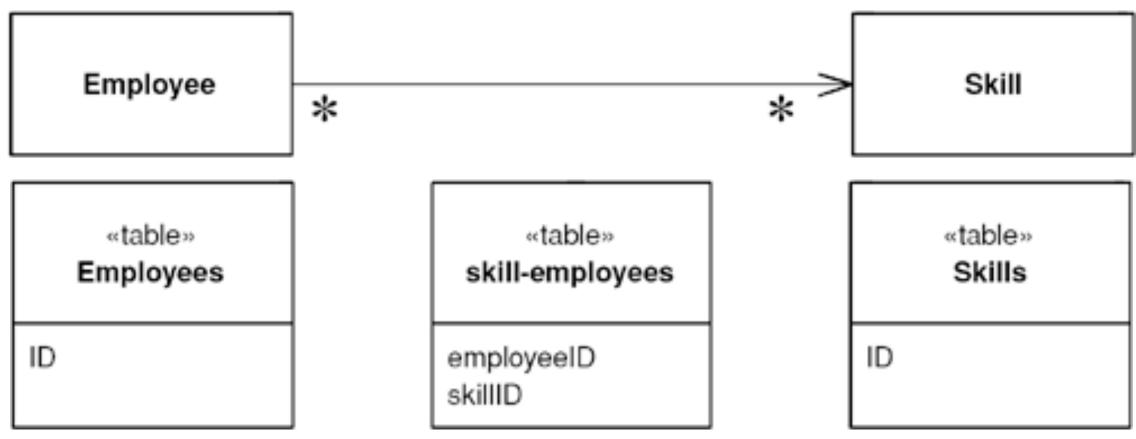






Association Table Mapping

• "Saves an association as a table with foreign keys to the tables that are linked by the association."







Many to Many



Many to Many

- This is the many to many association
- Allows for reuse of common data
 - category, skills, zip codes
 - Anything that multiple objects might be associated to and associated with
- Solved with an extra table





- Use a third 'link table' to store the association
- The table will typically only have the two foreign keys
 - You can store data specific to the association





- The link table will not normally have a corresponding data model object
- Can be loaded with
 - two queries (can be hard to coordinate)
 - using a join



When to Use It



When to Use It

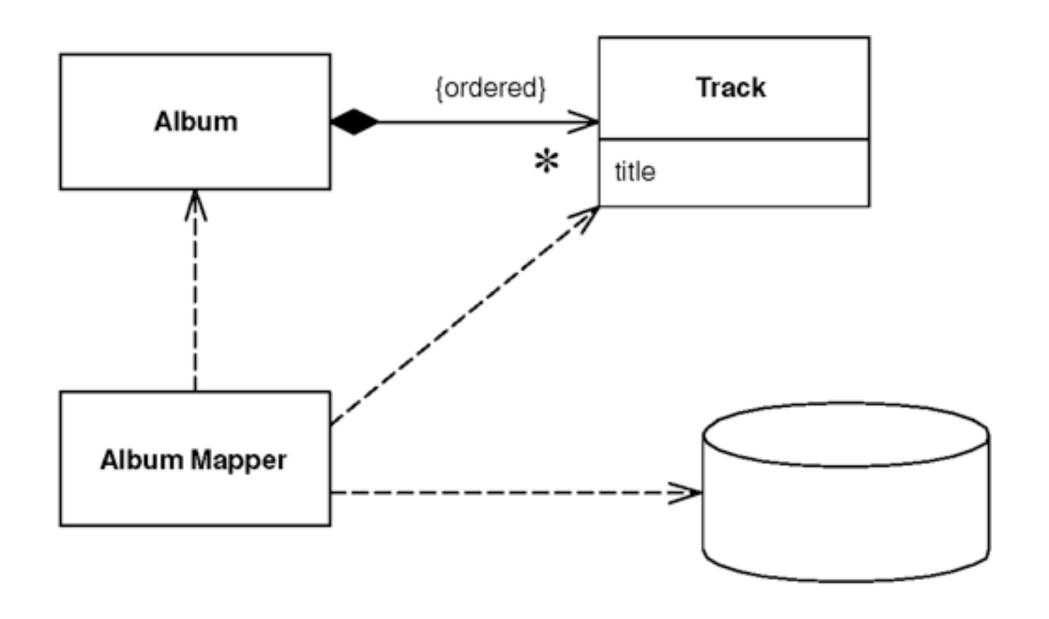
- When you have a many-to-many relationship to map
- When you have extra information to store
 - works fine, but you need the corresponding domain object

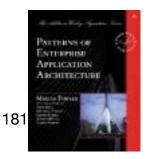


Example (from book)



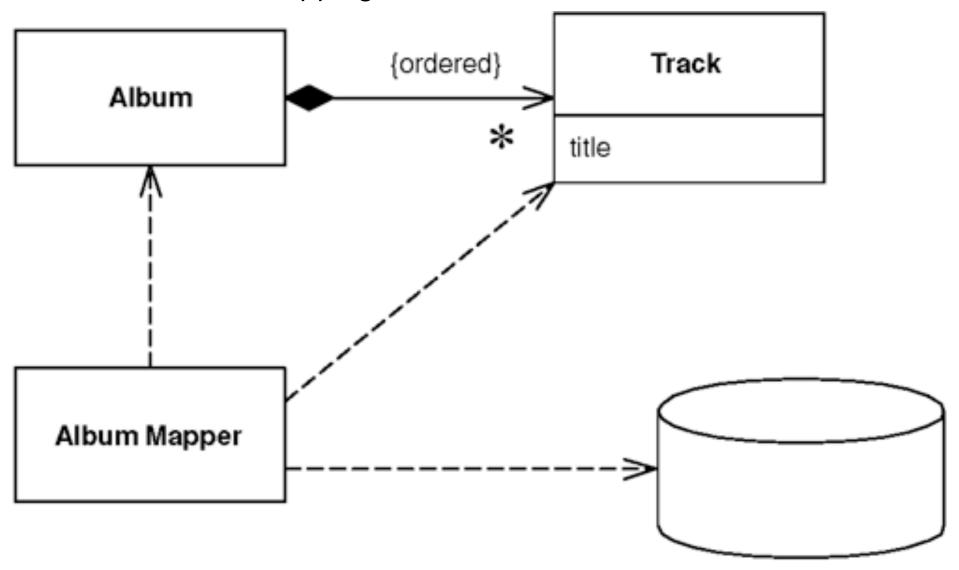


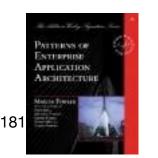






• "Has one class perform the database mapping for a child class."









- For objects that are naturally part of their parent object
 - but they are a collection, so they have to be stored on a separate table





- A dependent class relies on its owner
 - Every dependent must have an owner
 - And can not be loaded unless it is in the context of its owner





- Active Record, Row Data Gateway
 - There is no DB logic in the domain object
- Data Mapper
 - There is no mapper for the dependent object
- The dependent object can be lazy loaded



Identity



Identity

- The dependent object does not have an identity field
- It will be identified by the same ID as its owning object



Changes



Changes

- Changes to the dependent object
 - require you to save the owning object even if it hasn't changed on its own



When to Use It



When to Use It

- When you have an object (or collection) that only exists in the context of another object
- I like to use this when a record is large (say 30 fields) but not all of them are needed in all situations
 - Push to separate table & lazy load



Example (Book)





Employment

ID

person: person

period: DateRange

salary: Money

∝table» Employments

ID: int

personID: int

start: date

end:date

salaryAmount: decimal

salaryCurrency: char





• "Maps an object into several fields of another object's table."

Employment

ID

person: person

period: DateRange

salary: Money

«table» Employments

ID: int

personID: int

start: date

end:date

salaryAmount: decimal

salaryCurrency: char



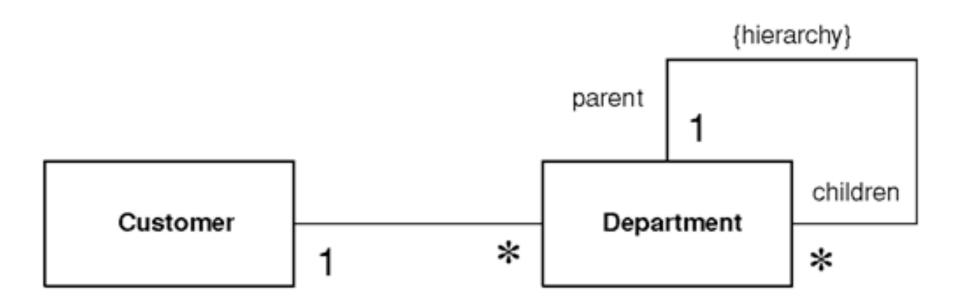


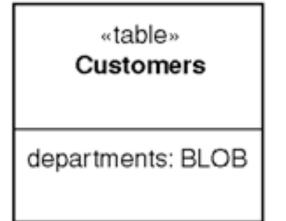


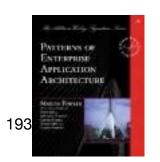
• Really no different than Dependent Mapping except it is explicitly a single dependent object, rather than a collection





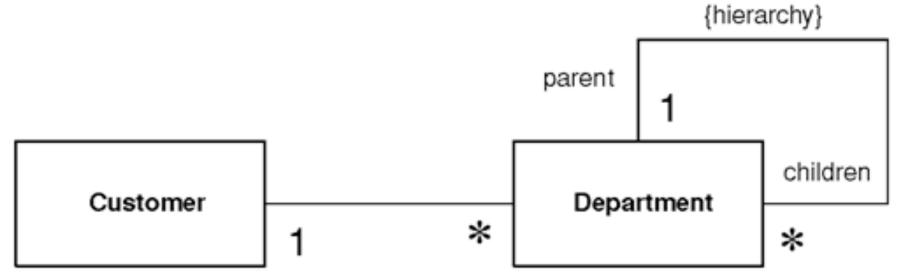


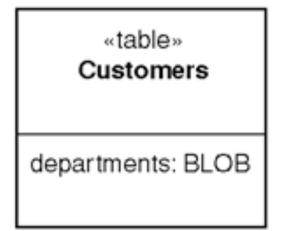






• "Saves a graph of objects by serializing them into a single large object (LOB), which it stores in a database field.









LOB



LOB

- Databases now have to "large object columns"
 - BLOB Binary Large OBject
 - CLOB Character Large OBject
- In HSQLDB
 - LONGVARBINARY, LONGVARCHAR





- Allows serialized objects to be persisted
 - i.e. the object doesn't have to be broken down into table / fields
 - Good for objects that you don't have to search
 - or they have a dynamic structure over time





- Two main types of serialization
 - Binary BLOB
 - Text/XML CLOB



Blob



Blob

- The Good:
 - Easy to program (Java binary serialization isn't so bad)
- The Bad:
 - DB must support BLOBs
 - · Can force lengthy and complicated conversions when the software is upgraded



CLOB



CLOB

- Can be platform independent
 - XML that can be written / read by multiple systems
 - Can be read by humans (DBAs, Developers)



CLOB



CLOB

- Downsides to XML
 - Parsing XML can be time consuming
 - XML also takes up much more space than the binary version of an object



Identity Issues



Identity Issues

- Since data in the LOB column has a singular ID (for the whole serialized list)
 - What happens if the same customer is saved in two different rows?
 - Is this a problem?
 - How do we identity the correct one?



Duplication



Duplication

- We can avoid this by limiting (or eliminating) duplication
 - Try to store data in the LOB column that pertains only to that particular record and doesn't overlap with another





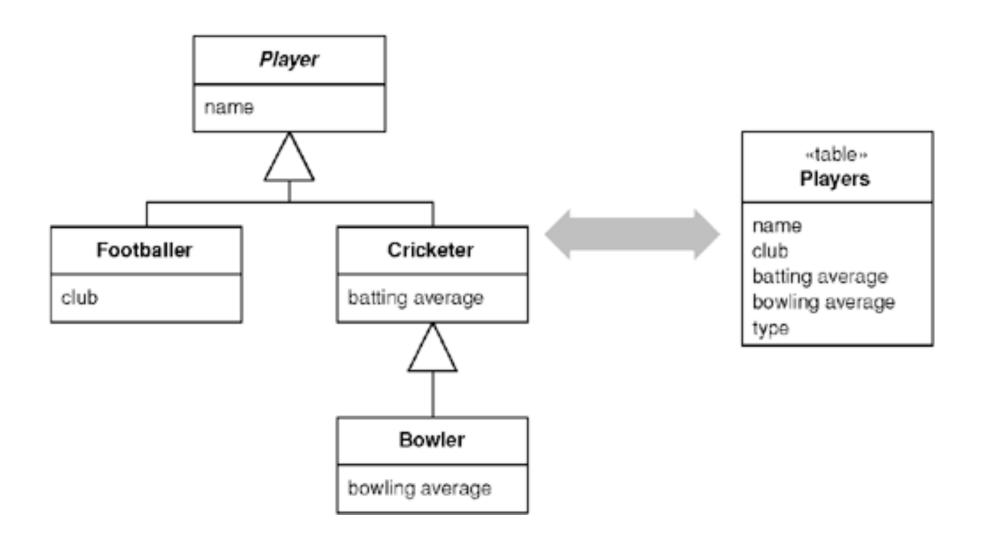
- A good way to store objects that will never be queried on
- When you can't alter the database structure
 - I've actually seen this



Example



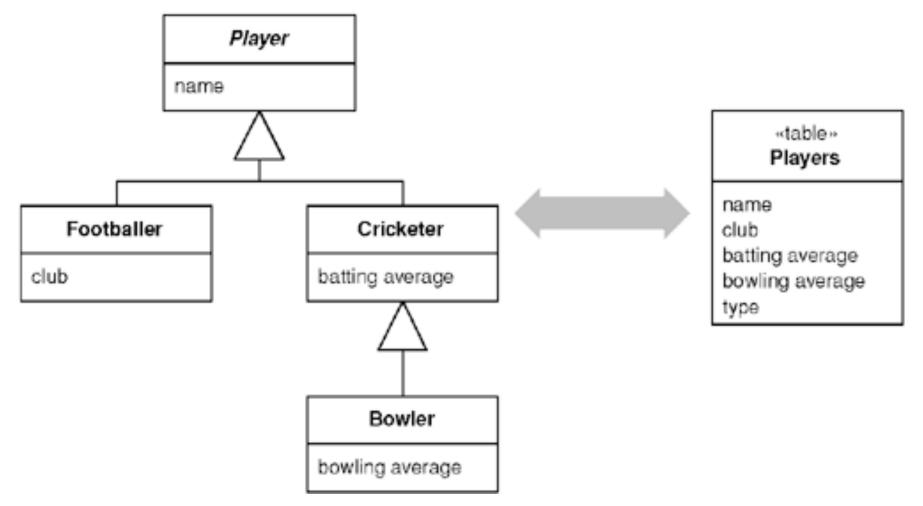








• "Represents an inheritance hierarchy f classes as a single table that has columns for all the fields of the various classes."









- As we noted earlier RDBMs don't support inheritance
- In many instances, it is nice to have inheritance in our object model
 - The resolution of this is in the mapping (one of the next 4 patterns)





- One table contains all data
 - for all classes in the hierarchy
- Each object is stored in a single row
 - This leaves the possibility of empty columns in any given row





- As each row is read
 - Determine the correct concrete class to instantiate
 - Typically there will be a field indicating which type it is (type code or class name)



Naming



Naming

- Using the class name
 - creates a coupling between the DB & object model
 - can be used to instantiate the object (using reflection)





- The Good:
 - only a single database table
 - no joins
 - field level refactoring doesn't force a database change

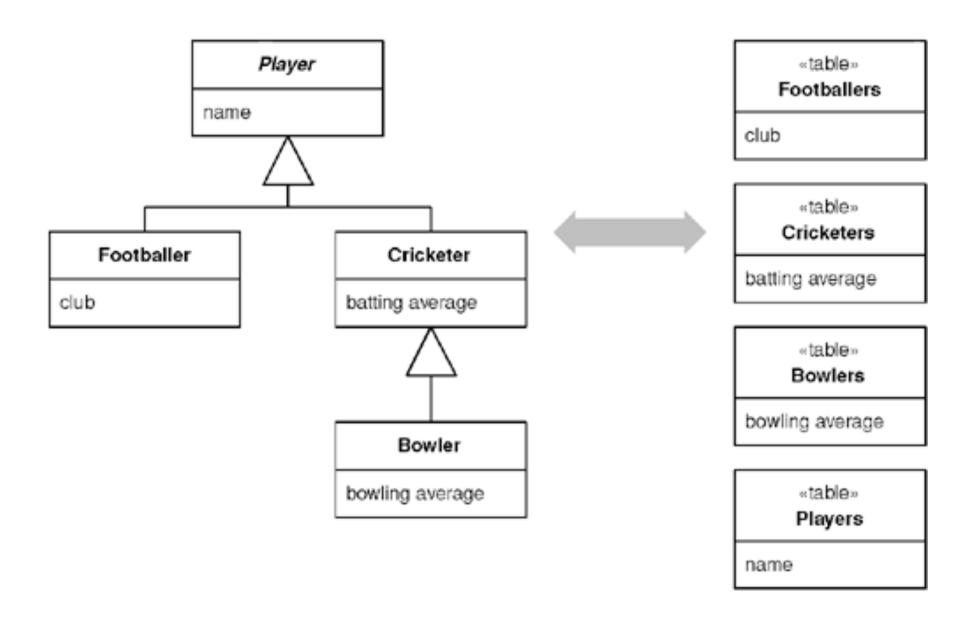


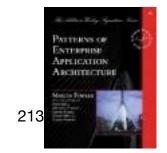


- The Bad:
 - seemingly random empty fields can be confusing
 - wasted space in database (most databases efficiently represent null values though)
 - the single table could be very large leads to contention
 - single namespace
 - can't use database to enforce required fields



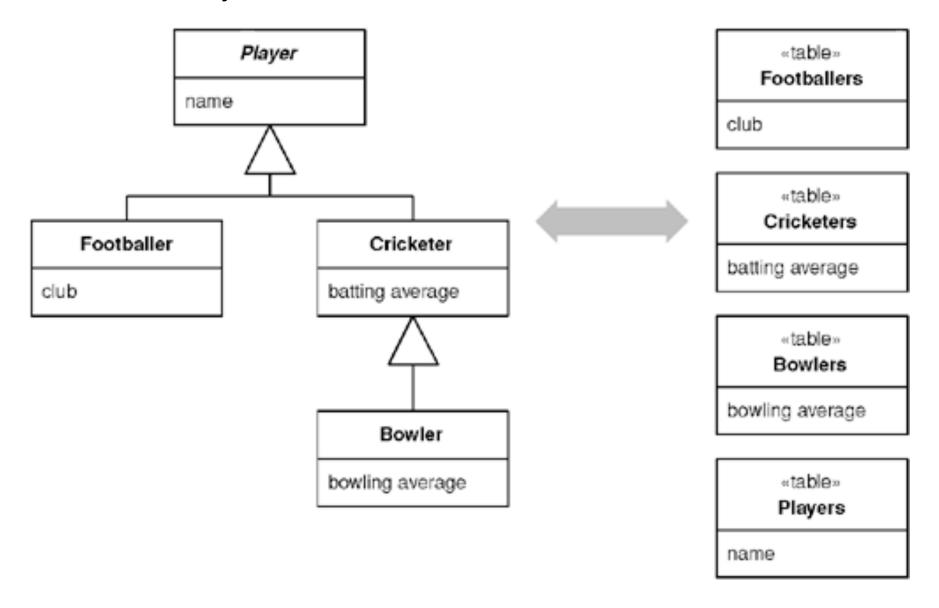


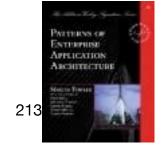






• "Represents an inheritance hierarchy of classes with one table for each class."







- Supports one table for every class in the hierarchy
 - Including the abstract classes





- Easy to understand with the 1 to 1 correspondence of table/class column/field
- You have to know how to join the correct values in
 - Use a shared primary key
 - Each key must be unique across ALL tables (football / baseball player can't both have an ID of 1)





- Need to do a complex join to get all of the data back
 - but if the hierarchy is complex enough, you might not be able to join everything together in one call
- Plus you don't know where to start reading





- The Good:
 - All columns are meaningful on all tables
 - Easy to see the relationships between the object model and database tables





- The bad:
 - need to read from multiple tables (joins)
 - Refactoring of fields causes changes to multiple tables
 - Ad hoc queries become difficult to write



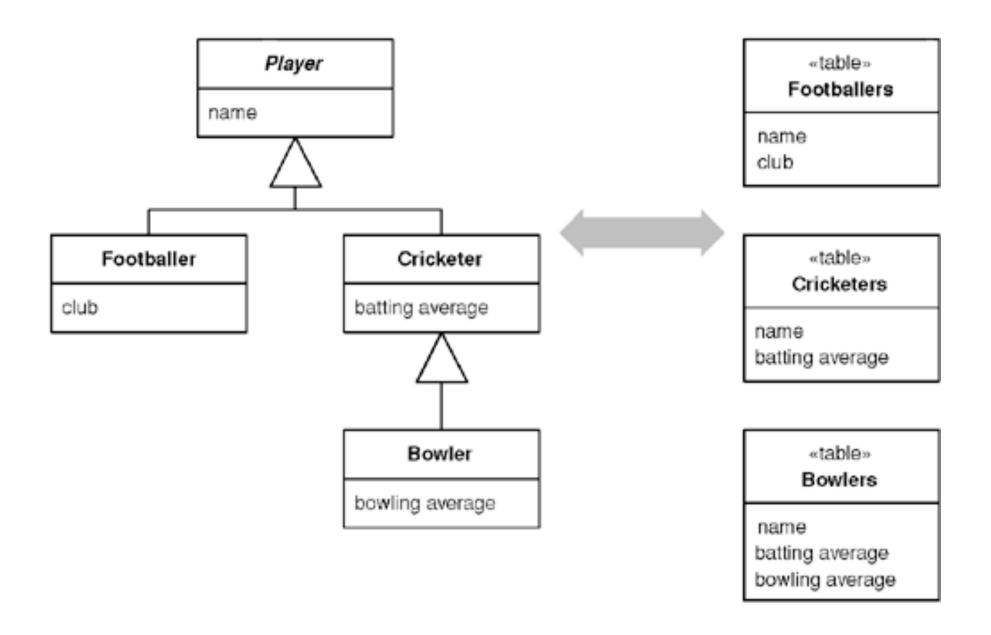
Example



Concrete Table Inheritance



Concrete Table Inheritance

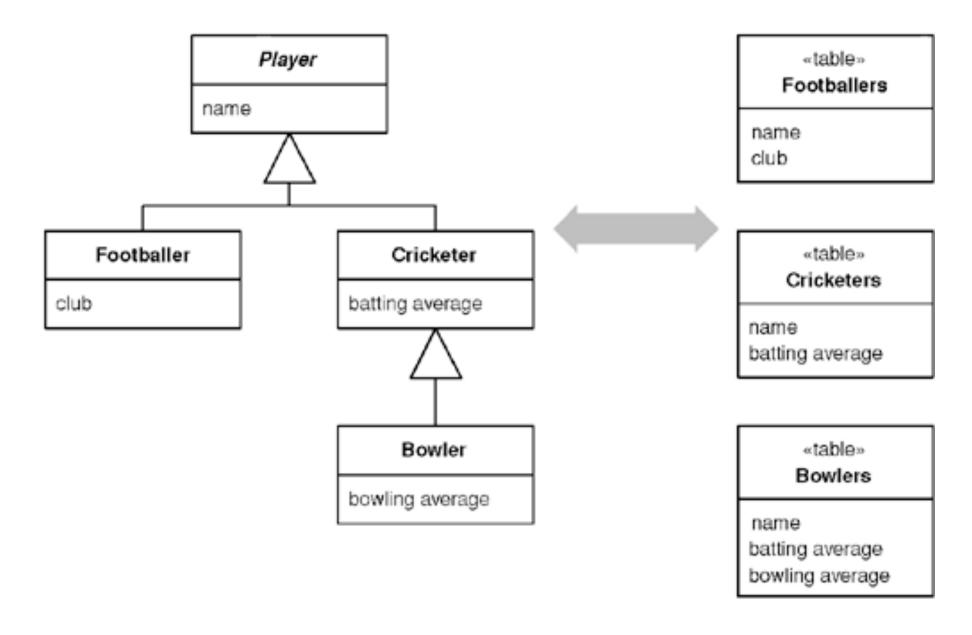


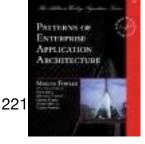




Concrete Table Inheritance

• "Represents an inheritance hierarchy of classes with one table per concrete class in the hierarchy."







- One table for each concrete class in the object hierarchy
- Each table contains the necessary columns for all fields in that silo of the hierarchy
 - some fields appear on multiple tables





- Keys are important
 - They must still be unique across all tables in the hierarchy





- Finding the correct object by ID is difficult.
- You need to search all of the appropriate tables





- The Good:
 - · each table contains all of the necessary data to stand on its own
 - no joins to load
 - · tables can balance the load
 - a somewhat empty argument





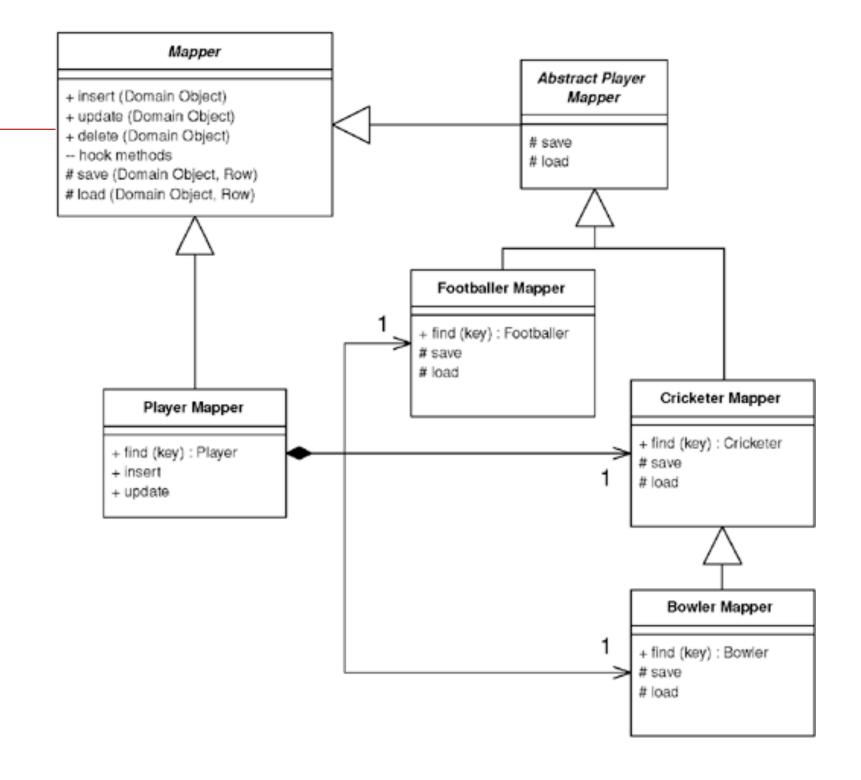
- The bad
 - must have unique IDs across tables
 - hard to enforce relationships with abstract classes
 - refactoring may require database changes
 - a change at the superclass requires changes to all tables
 - find requires multiple queries



Example



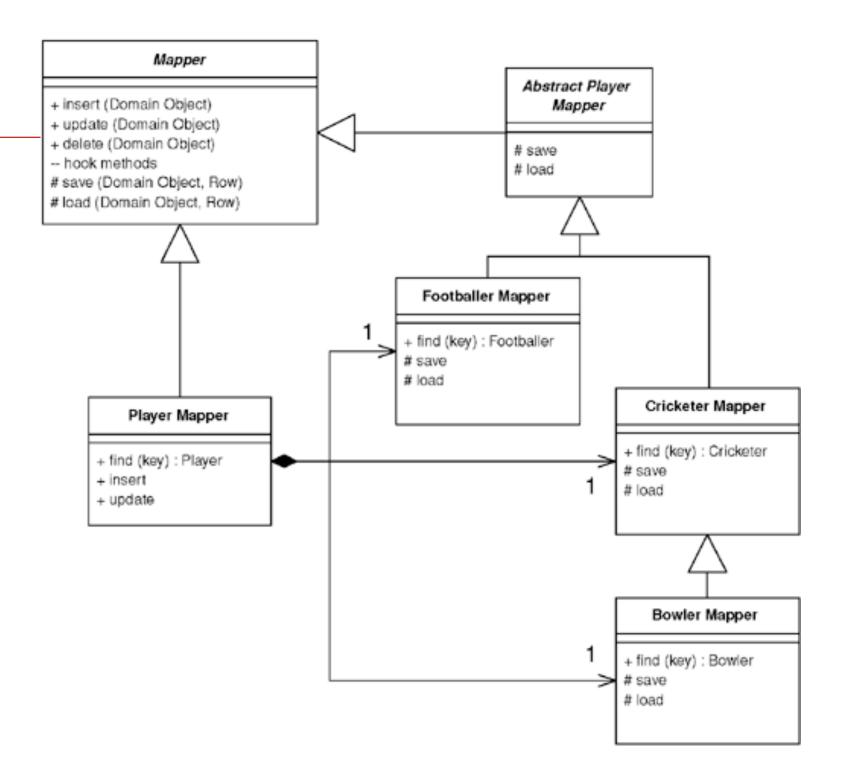








• "A structure to organize database mappers that handle inheritance hierarchies."









- We want to minimize the amount of code used to load/save hierarchies
- We also want to separate
 - abstract reusable portion of the mapper
 - concrete class specific portion of the mapper





- A hierarchy of finders corresponding to the domain object hierarchy
- Find methods
 - locate the appropriate database row and instantiate the correct objects
 - once you have the correct object it should be easier to save





Any time you have inheritance involved in the object hierarchy





Object-Relational Metadata Mapping Patterns (Patterns of Enterprise Application Architecture - Chapter 13)

Mike Helmick Large Scale Software Engineering - Spring 2014 University of Cincinnati

Patterns (Ch 13)

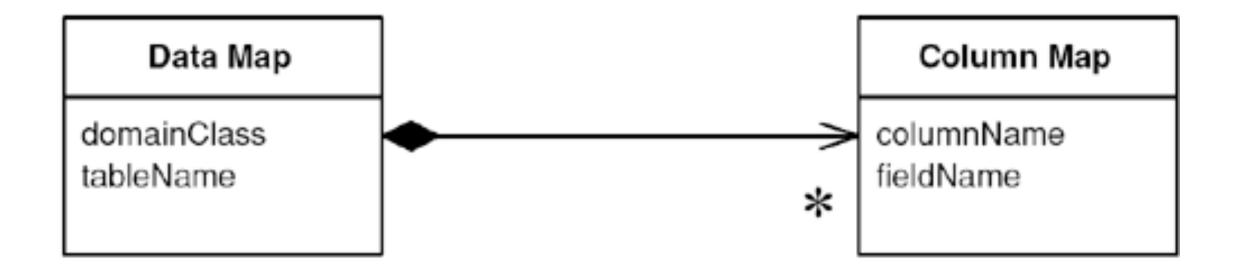


Patterns (Ch 13)

- Metadata Mapping
- Query Object
- Repository



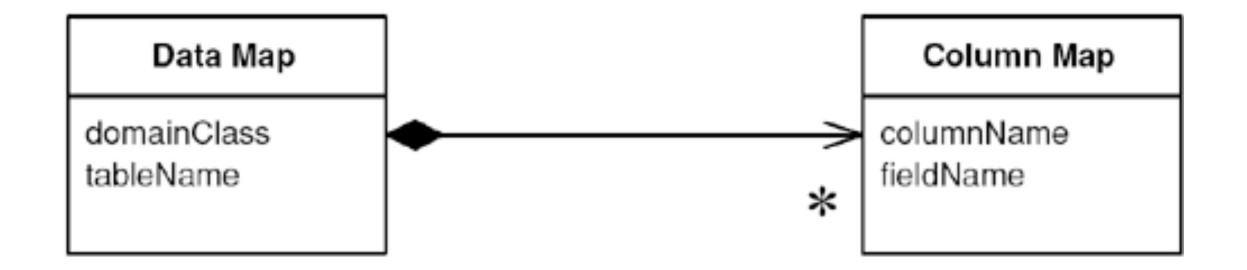








• "Holds details of object-relational mapping in metadata."









- So far we've seen a lot of hard coding of mapping into individual classes
- Not such a great approach is it?





- On a basic level
 - metadata is turned into code at runtime
- Two basic architecture approaches this is a big decision
 - code generation
 - reflective programming



Code Generation



Code Generation

- Write a program that reads the metadata and outputs source code
- The classes look like someone coded them
- Usually done as part of the systems build
 - No reason to check these in to a source code control system
- Code should never be edited



JIT Code Generation



JIT Code Generation

- This technique is not in the book
- It is possible to generate code at runtime
 - In Java this is throw dynamic subclassing and byte code generation
 - I think this is a very valuable approach, but requires coding to interfaces
 - Better suited to dynamic languages (Ruby)





- Treat both the methods and fields as data
- Find a field called name
 - ask the object if it has a method called setName
 - Invoke the method dynamically





- Author says he usually recommends against it
 - hard to debug
 - slow (when comparing to a standard function call)
- But acknowledges that it can be appropriate





- My opinion:
 - I am actually a big fan of the reflective properties of modern OO languages
 - They can seriously simplify generic code making it easier to maintain
 - I will almost always sacrifice a little runtime speed for maintenance speedup



Code Generation



Code Generation

- Changes to the mapping layer
 - require regeneration and full build
 - or just restart server if it is dynamic generation



Debugging



Debugging

- Both approaches are somewhat difficult to debug
- With generated source code you can step into the methods (if you keep the source code around that long)
- Less experienced developers will have trouble debugging a reflection based mapping layer





- Needs to be stored somewhere
- XML
 - Hierarchical format
 - can have a table tag, field tags inside
 - Parsers / editors / languages (XQuery) already exist





- Some other format you have to write your own parsers
- Can actually be stored in code
 - As compiled source
 - As annotations (Java 5 has these now)





- Store it in the database
 - As a separate mapping table
 - Or just read the table structure
 - JDBC metadata API





- Speed of accessing / loading the metadata can be ignored
 - Even if it takes a while
 - Since it should only be loaded once and kept around
 - As generated code or as generated reflective mappings



Complexity



Complexity

- You probably won't be able to guess at everything you need when constructing this layer
- Take an evolutionary approach to development



Special Cases



Special Cases

- For most of what we do (CRUD type operations)
 - Some complex mappings can't be expressed with the same metadata
 - Or the effort to code the mapping generator is too great
 - You can always sub-class generated code to modify behavior
 - Handle these situations as they arise





- Metadata mapping can greatly speed up development
- Developing the mapping framework is quite complex
 - Takes a good amount of time requires (generally) a senior developer / architect needs to be right





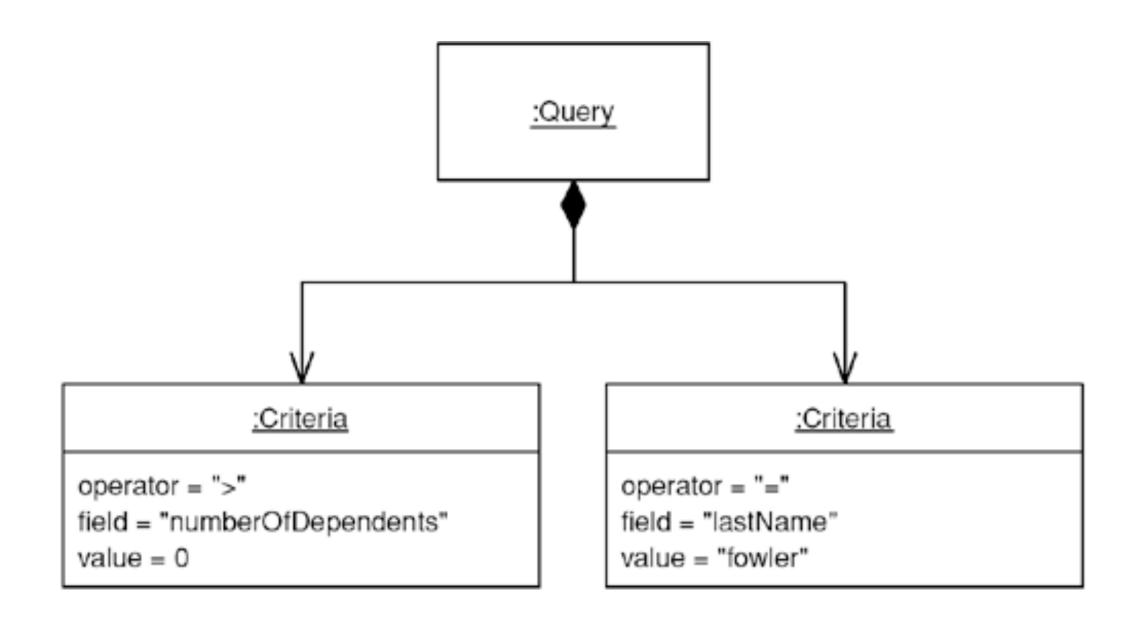
- In general:
 - I think this is the best approach
 - Many commercial / OSS tools exist for a variety of programming languages



Example



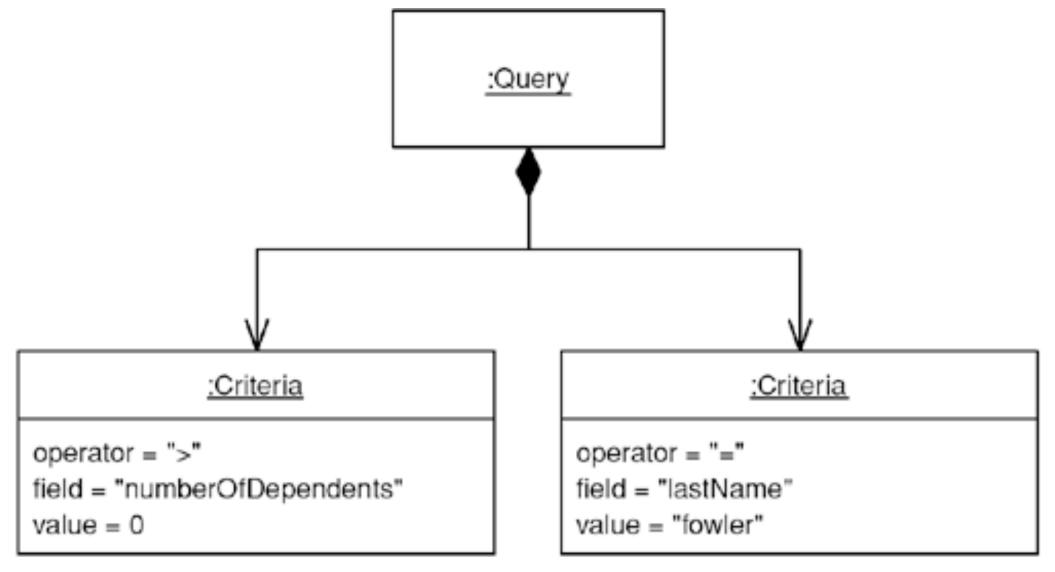








• "An object that represents a database query."









- On a given project there is a good chance that not all of the developers will know (or know well) SQL
- Queries are created by referencing fields and objects



How It Works



- Turns in-memory objects into an SQL query
- A simple implementation will usually work to start with
 - and can be expanded upon as new features are required





- Queries are formulated in the programming language that you are using
 - Often times the use of the field name is exactly the same as using the database column name
 - This can help to limit the impact of changes in one or the other
 - Assuming a good mapping layer





- Can help isolate the developers from the differences in SQL from one database to another
- Hibernate / HQL does this



Variation



Variation

Use a partially populated domain object to infer a query





- Doesn't work with a hand coded data source layer
- works well with the combination of
 - Domain Model, Data Mapper, Metadata Mapping
 - Domain Model, Active Record, Metadata Mapping





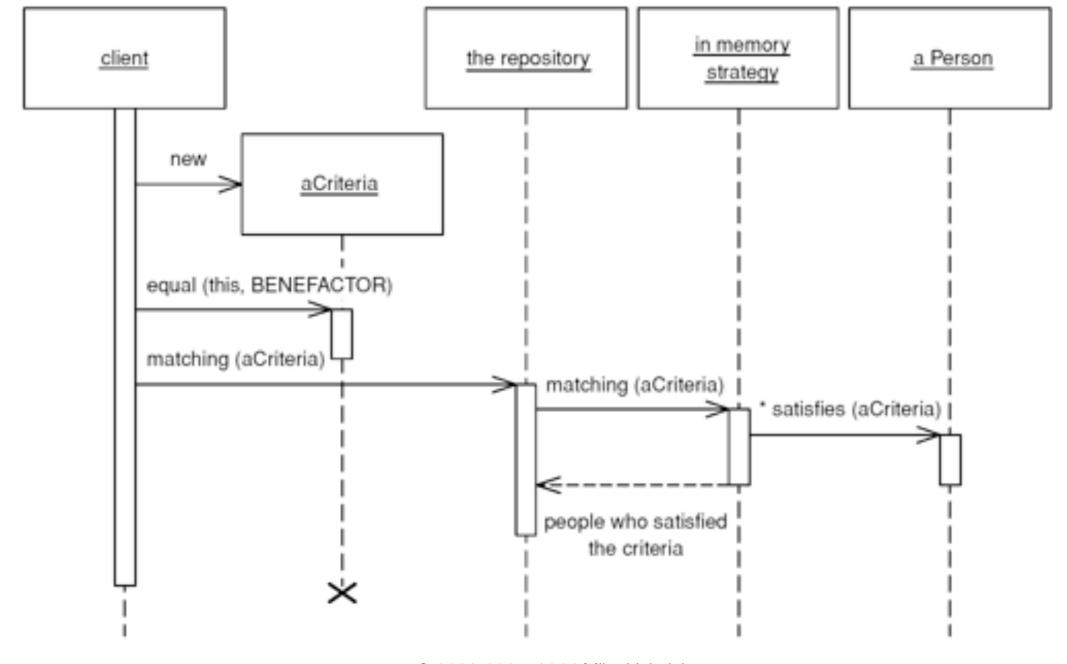
- Developers aren't too familiar with SQL
- You have a tool that can be purchased to accomplish this goal
 - These are very difficult to construct

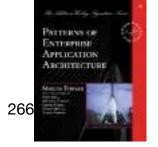


Example



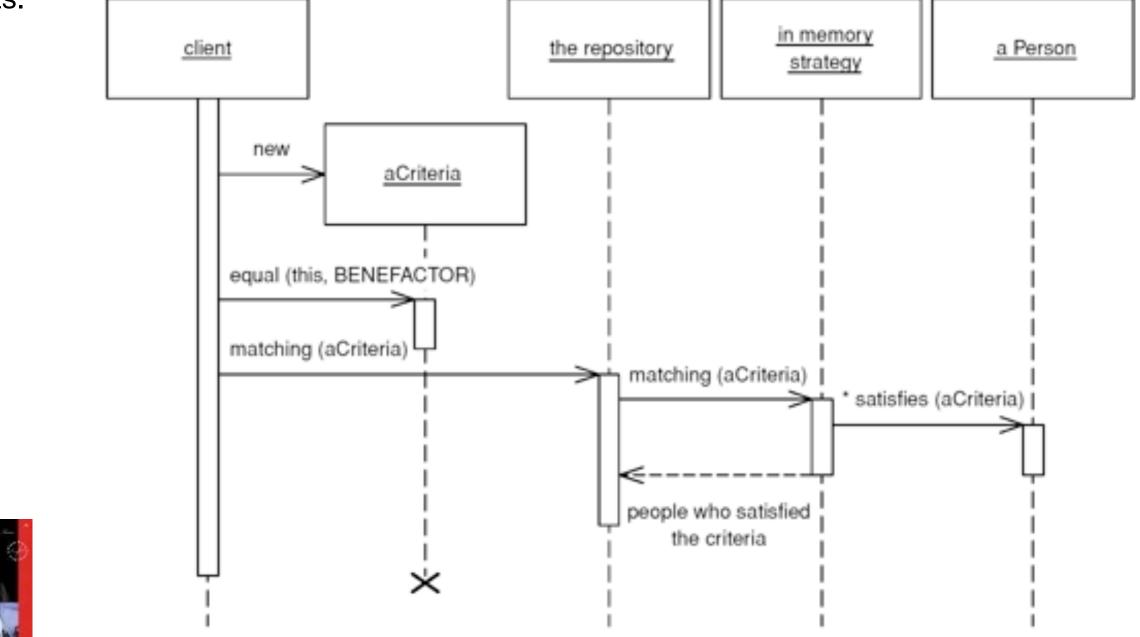






• "Mediates between the domain and data mapping layers using a collection-like interface for accessing domain

objects."





APPLICATION



- Another layer that abstracts out querying functionality
- Eliminates coupling between query building and the data access layer





- Works well with the query object
 - And while difficult to build
 - Not a far leap if you are already using Query Object





- Create a criteria based query
- The repository then presents the database as if all objects are in memory and you are just asking for a subset
- Eliminates the need to specialized finders (Java)





Uses Metadata Mapping and Query Object to power its functionality





- SQL is totally removed
- All queries are done based on the in-memory objects

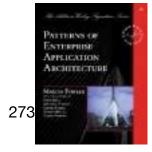




- Multiple data sources
 - When your application is pulling data form more than one database
 - This can encapsulate this detail
- Can also encapsulate a simulated database (testing)



Example





Example

- From the book
 - Page 325

