

# Concurrency, Session State, Distribution

---

Large Scale  
Software Engineering  
Mike Helmick  
University of Cincinnati  
Spring 2014





# Patterns of Enterprise Application Architecture

---

- Now, we shift back to software architecture and specific architecture patterns
- Starting with Chapters 5, 6, and 7 (this presentation)



# Concurrency

---

# Concurrency

---

- A very difficult subject
- Pervasive in enterprise applications
- Many users, many processes, many threads

# Concurrency

---

# Concurrency

---

- Something that is going to happen
- Something that is often ignored (wrongly)

# Why Ignored?

---

# Why Ignored?

---

- Because we have transaction managers
  - Pieces of software that can manage aspects of concurrency for us
  - i.e. Do everything in a managed transaction, and most things work out



# But...

---

# But...

---

- This doesn't always work well with items that span multiple database transactions
- This is called offline concurrency

# In the server

---

# In the server

---

- We also have concurrency in the application server
- Multiple threads of execution will be in the same code at the same time

# Concurrency Problems

---

# Concurrency Problems

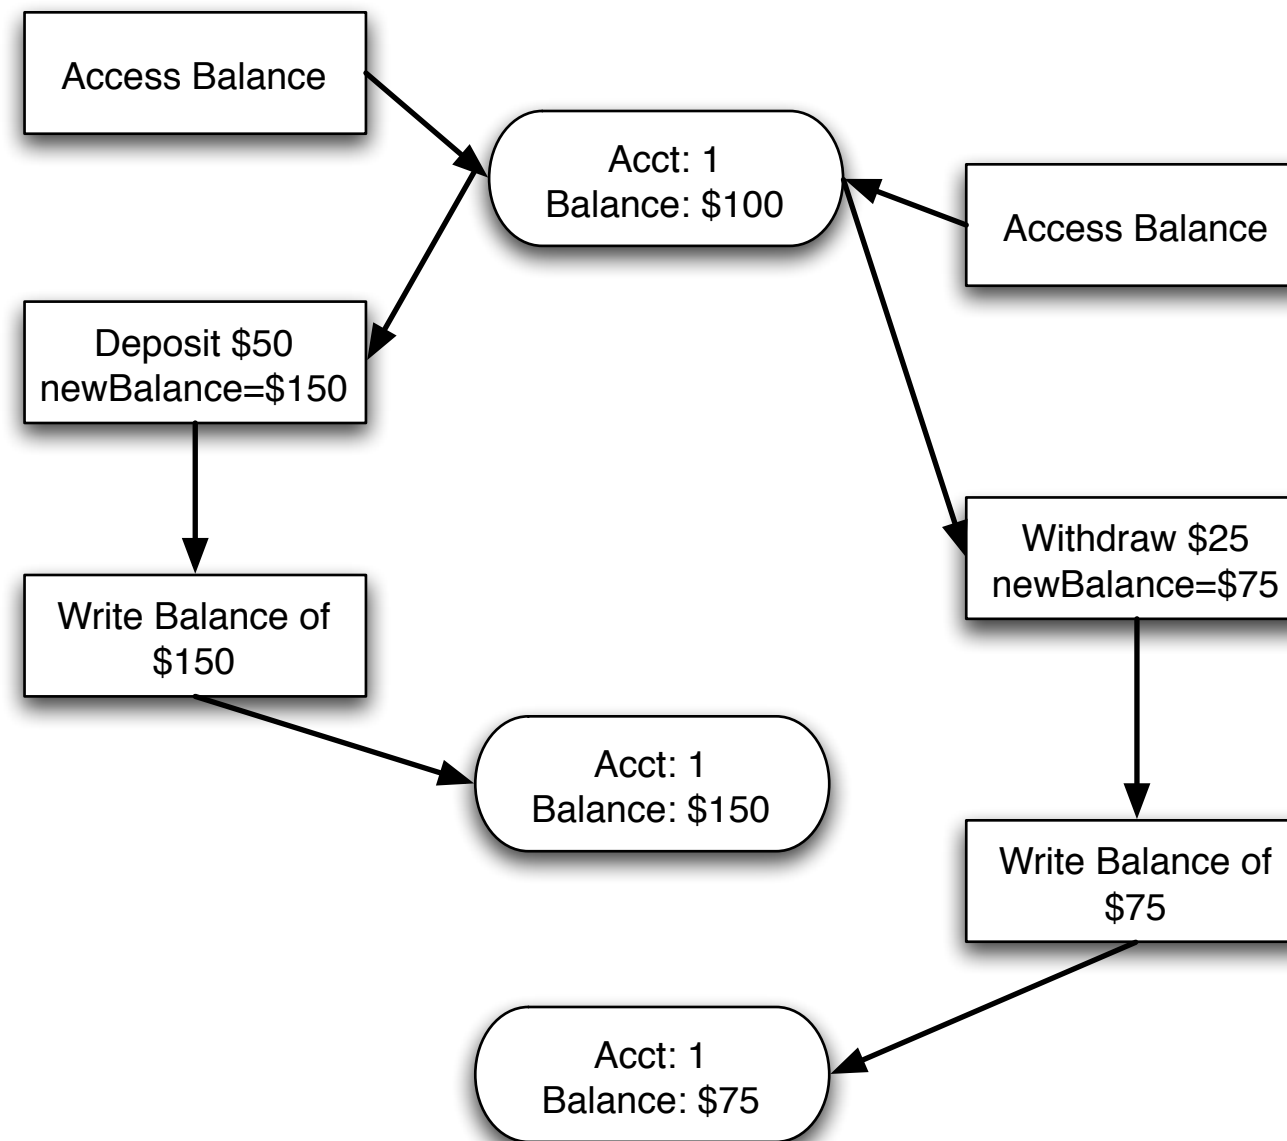
---

- Lost Updates



# Concurrency Problems

- Lost Updates



# Inconsistent Read

---

# Inconsistent Read

---

- Occurs when you read two (or more) separate pieces of information
  - Read first
  - First changes
  - Read second
- If the two pieces of data go together - they are now inconsistent

# Correctness

---

# Correctness

---

- What we want is correctness
- Data that is always in a correct and consistent state

# Liveness

---



# Liveness

---

- How much can we do at the same time
- May have to sacrifice some correctness for liveness

# Execution Contexts

---

# Execution Contexts

---

- Generally we talk about where something executes in some context
  - request
  - session

# Request

---

# Request

---

- A single call into the system
- Work is done and a response is sent back
- Best if the client must wait for a response
  - not the case on the web
  - and the request appear linked to the user - but not to the server

# Session

---



# Session

---

- A series of requests over time between a client and a server
- Could be one or more requests
- Often login through logout

# Session

---

# Session

---

- With the web
- We use a session cookie to identify a user's activities over time
- The Servlet API has methods for this

# Process / Thread

---

# Process / Thread

---

- process
  - heavyweight execution context
- thread
  - lightweight execution context

# Requests

---



# Requests

---

- Threads give us multiple requests in a single process
  - with a shared memory space
  - shared memory can be a concurrency issue

# isolated threads

---

# isolated threads

---

- Some environments allow you to specify memory on a per-thread basis
- In Java
  - Thread local storage

# Contexts

---

# Contexts

---

- We have a problem in that execution contexts don't always line up correctly
- The same client might not talk to the same server for each of the requests in its session
- We can solve this in several ways

# Isolation

---

# Isolation

---

- A solution to multiple threads accessing the same data at the same time
- Partition data so that only one process can access a piece of data at a time
  - locks

# Isolation

---



# Isolation

---

- Reduces the change of data errors
- Create isolation zones within a program where operations are done safely

# Immutable

---

# Immutable

---

- Some data can be declared immutable
- This is a good idea if we're never going to change some reference data
  - A table containing dates

# Control

---

# Control

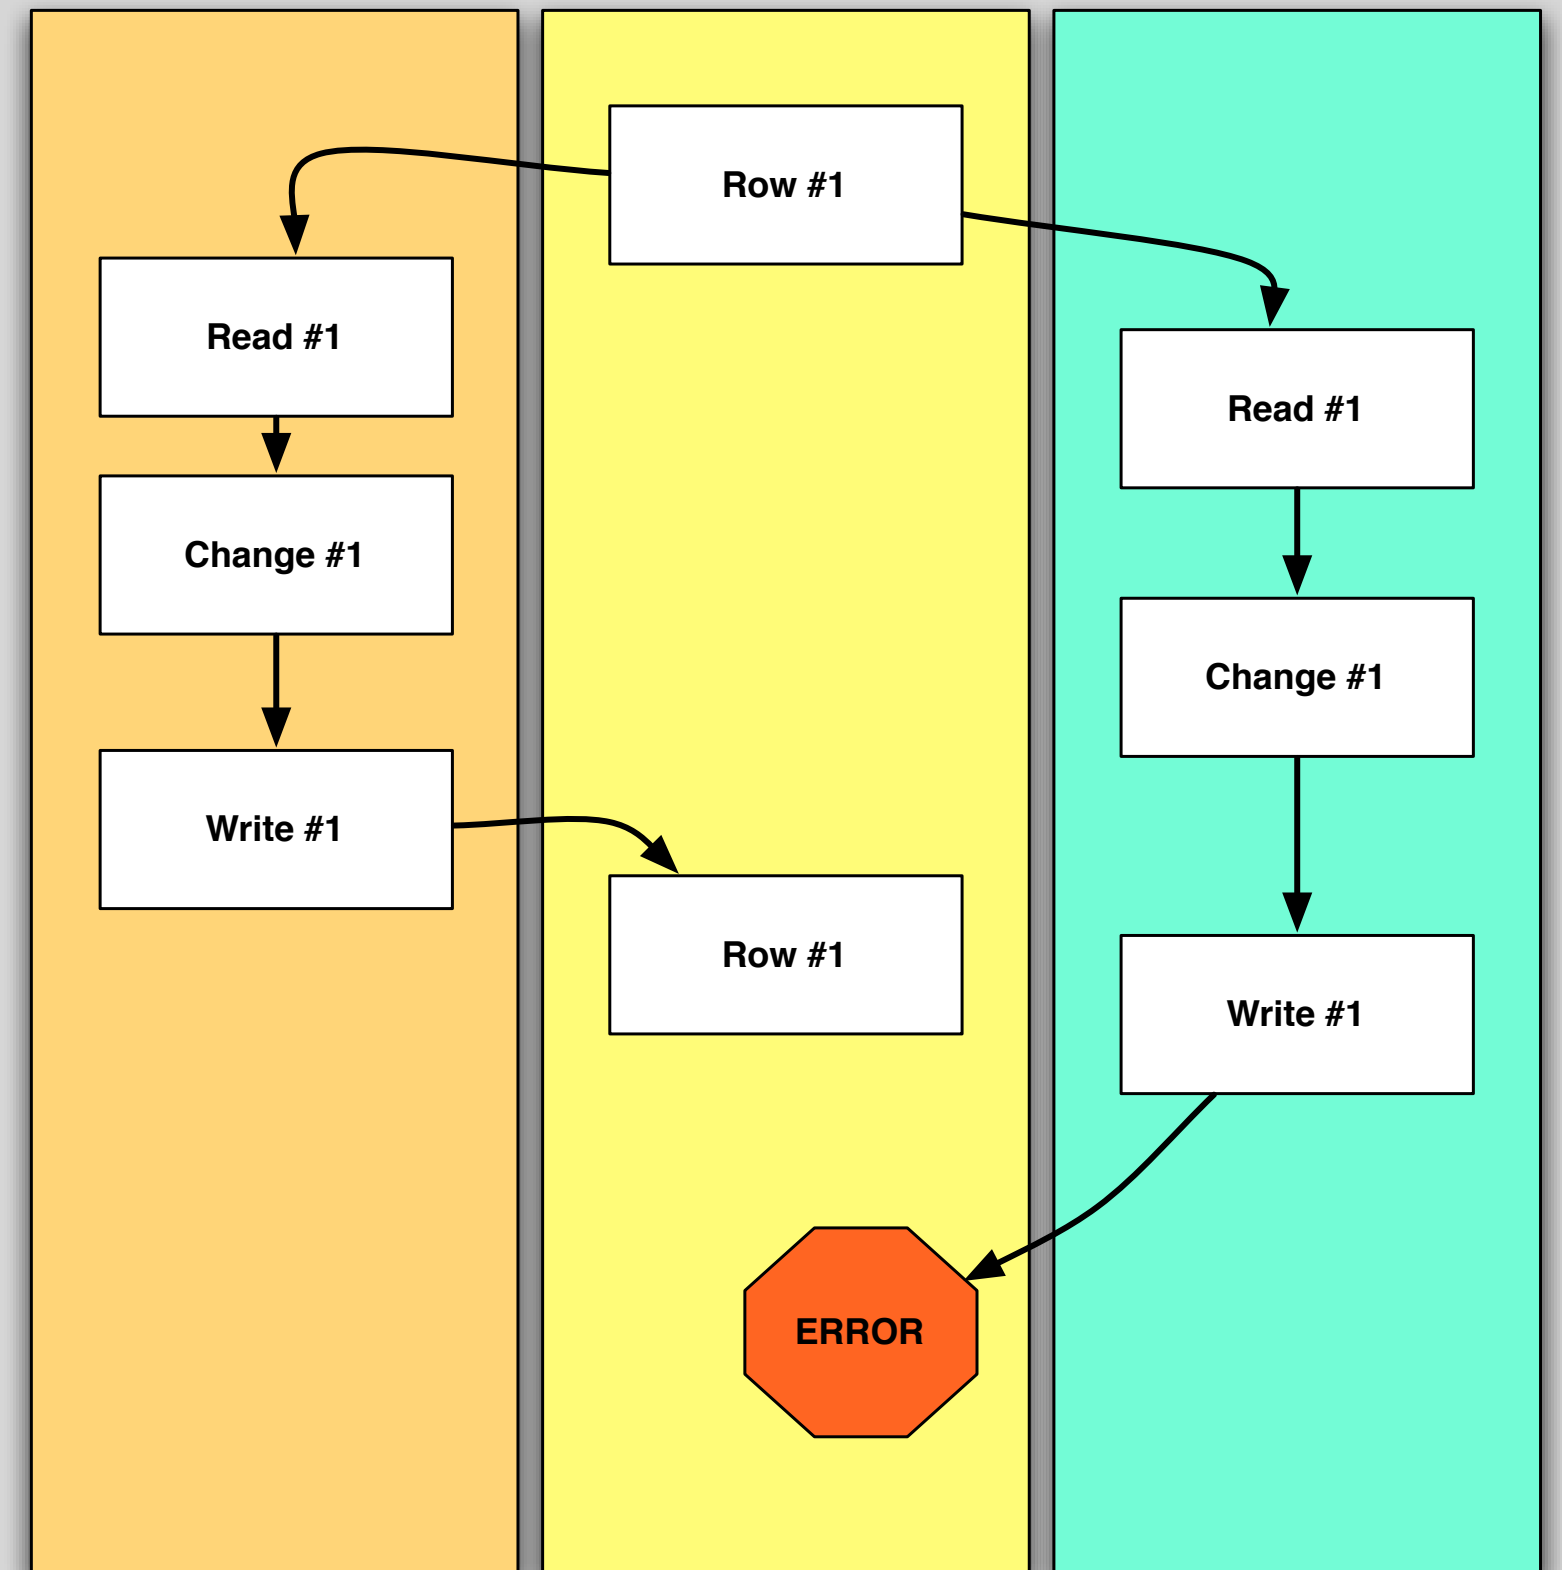
---

- In general - there are two types of concurrency control we use
  - optimistic locking
  - pessimistic locking

# Optimistic

---

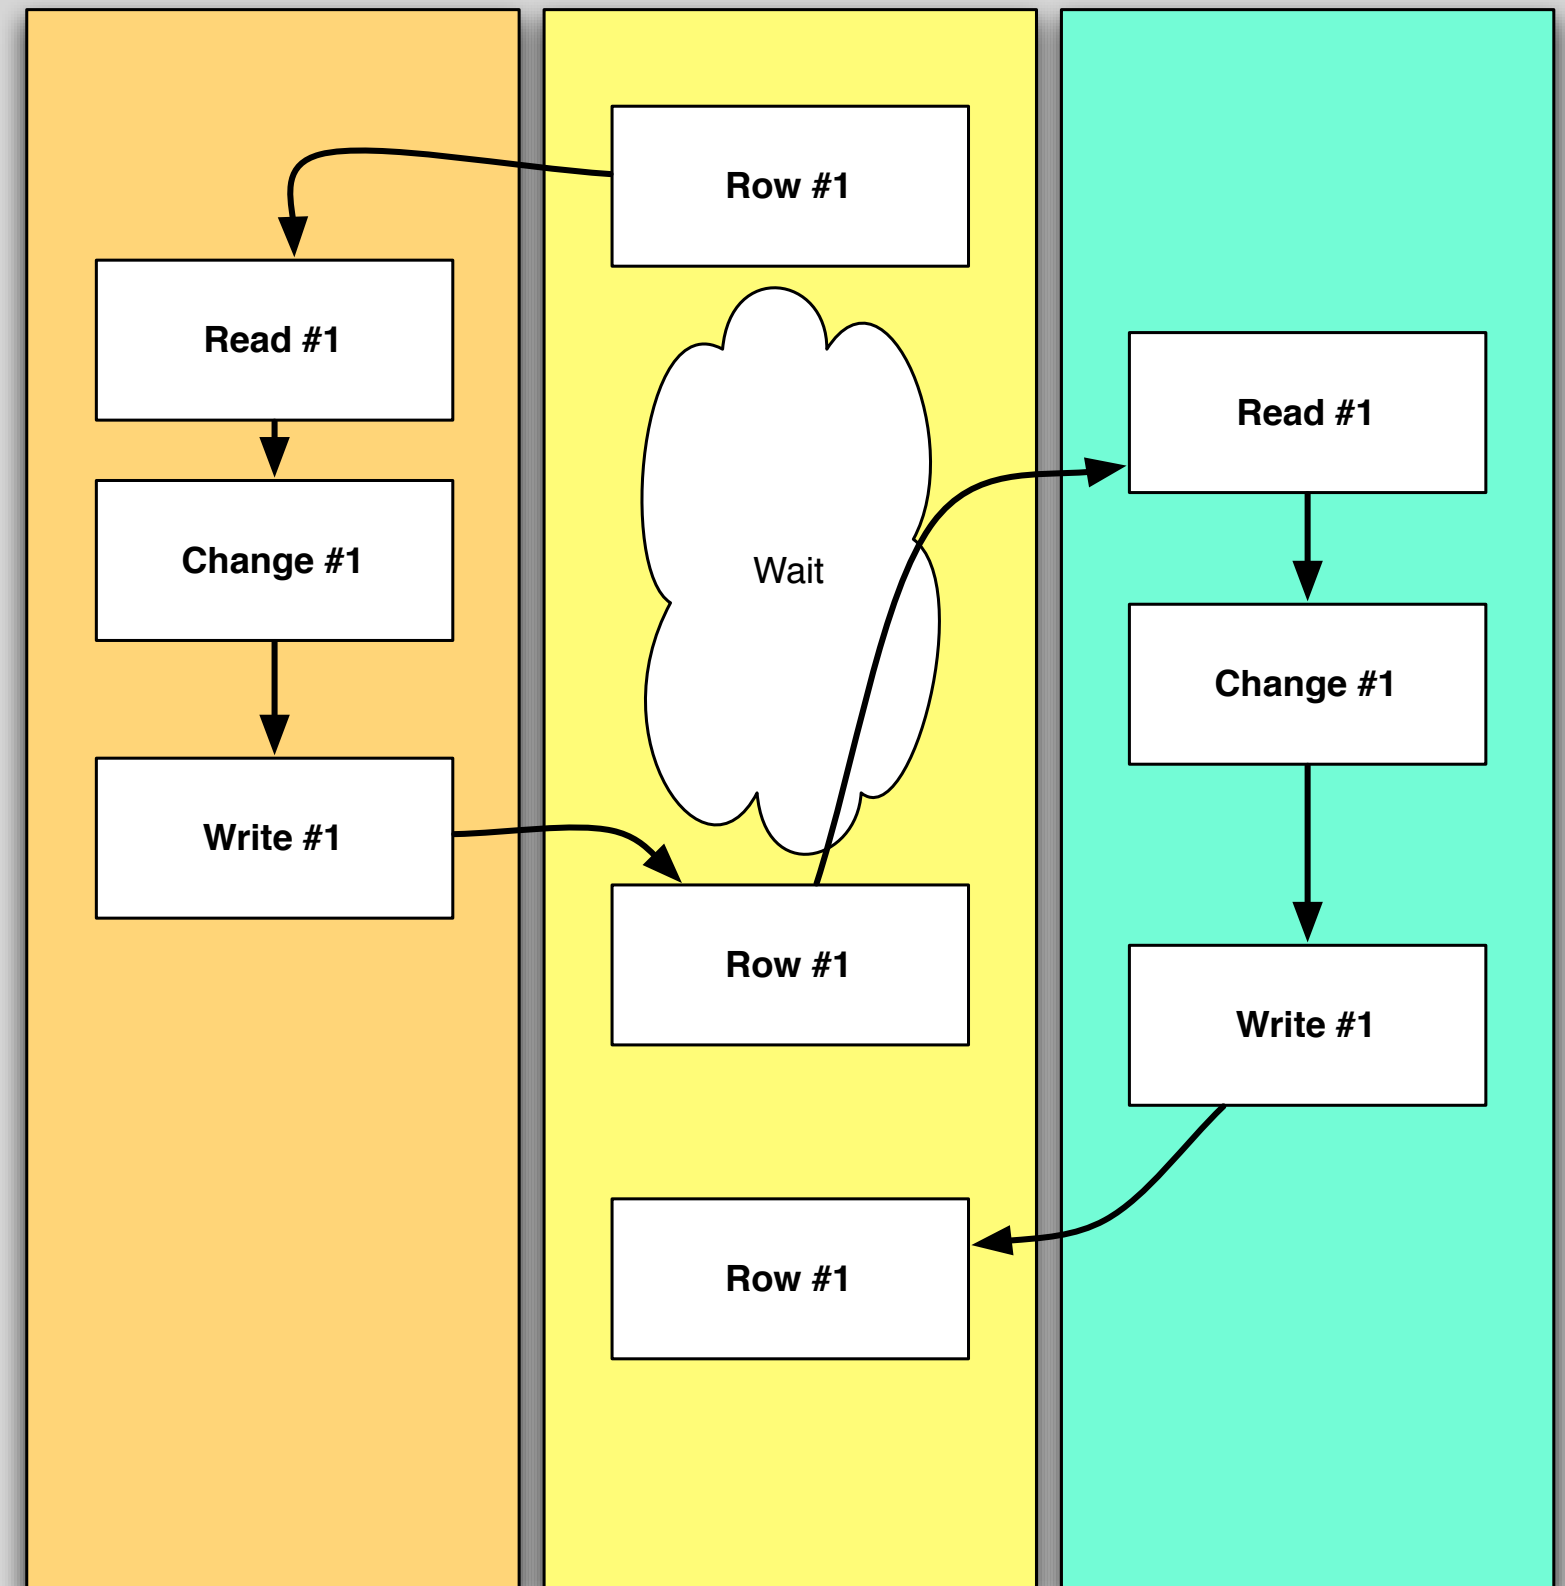
- Multiple copies of a record can be made
- but the first one to commit wins
  - Everyone else has to figure out what to do



# Pessimistic locking

---

- Who ever asks for the data first gets it
- All other requests are queued
  - They get the data when
    - it is committed
    - the transaction is aborted



both

---



# both

---

- Optimistic locking is used for conflict detection
- Pessimistic locking is about conflict prevention

either

---

# either

---

- Real systems can use either one
- For source code control systems
  - we tend to prefer optimistic locking
  - i.e. software development would halt if only one person could edit a file at a time

# choosing

---

# choosing

---

- pessimistic lock reduces concurrency
  - prevents data from being read while it is being edited
- supported natively by most databases

# choosing

---

# choosing

---

- optimistic lock
  - allows for greater concurrency
  - lock only happens during commit
  - causes merging problems - or transaction failures
  - not supported natively by most databases

# choosing

---



# choosing

---

- optimistic locking is usually the best choice
- but if data is really sensitive or can not be easily edited by multiple people at the same time - then pessimistic locking

# Deadlocks

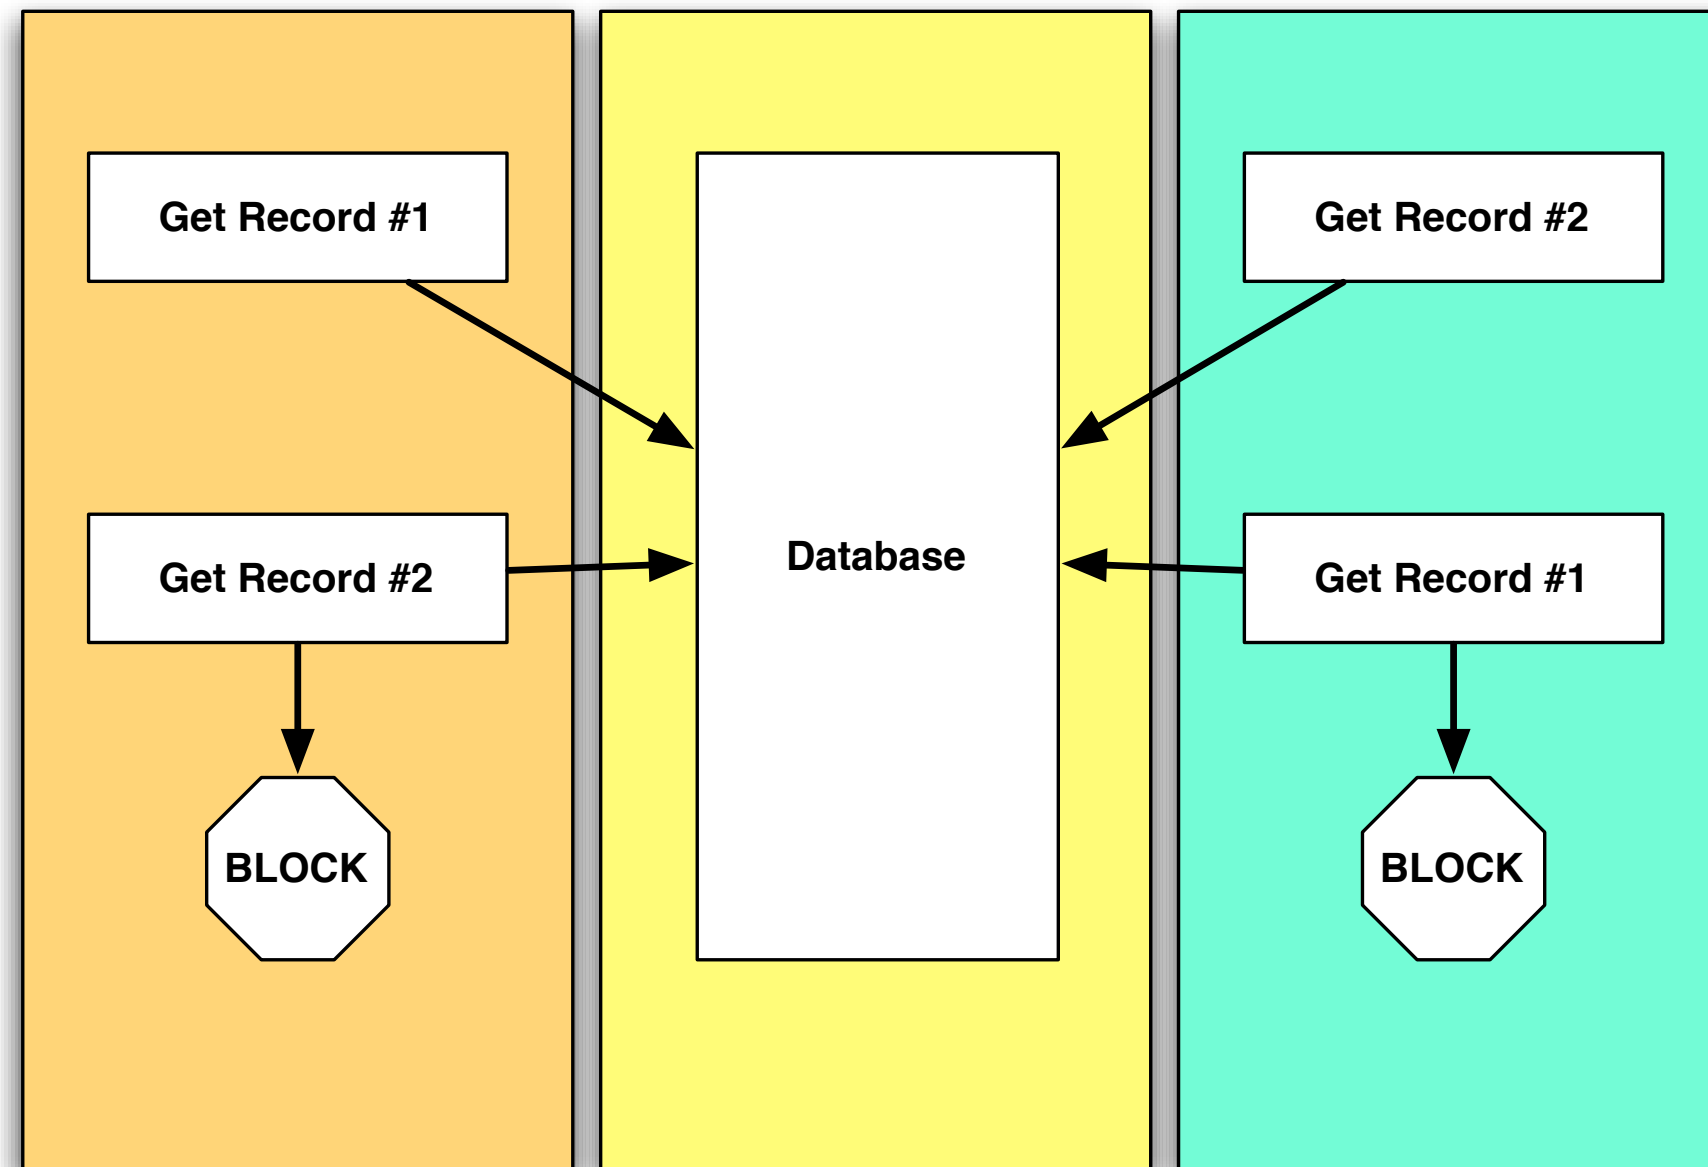
---

# Deadlocks

---

- A particular problem with pessimistic locking techniques
- Also called circular lock

# Deadlock



# Deadlocks

---

# Deadlocks

---

- Can be detected (see an Operating Systems book)
- We can attempt to prevent them

# Deadlock

---

# Deadlock

---

- occurs when
  - A user who already has locks - asks for more & won't release the ones they have



# Preventing Deadlock

---

# Preventing Deadlock

---

- We can force an order on how locks are requested
- This can avoid it

# Transactions

---

# Transactions

---

- A bounded sequence of work
  - defined starting point
  - defined ending point

# Transactions

---

# Transactions

---

- all participating resources are in a consistent state
  - when the transaction begins
  - when the transaction ends

# Transactions

---

# Transactions

---

- Are an “all-or-nothing”
- An ATM will not give you money without being certain that the balance of your account has been adjusted



# ACID

---

# ACID

---

- Atomicity
- Consistency
- Isolation
- Durability

# Atomicity

---



# Atomicity

---

- Each step in the transaction must complete successfully or the entire transaction will roll back



# Consistency

---



# Consistency

---

- “A system’s resources must be in a consistent, non-corrupt state at both the start and the completion of a transaction.”



# Isolation

---



# Isolation

---

- “The result of an individual transaction must not be visible to any other open transactions until that transaction commits successfully.”





# Durability

---



# Durability

---

- “Any result of a committed transaction must be made premanent. This translates to `must survive a crash of any sort.`”



# Transactional Resources

---

# Transactional Resources

---

- A transactional resource is anything that can use transactions to control concurrency

# Duration

---

# Duration

---

- We generally want transactions to be short
- If transactions span multiple requests, we call this a long transaction
- The common approach is to have one transaction for each request



# Duration

---

# Duration

---

- It is possible to read everything
- then open the transaction for updates
- this is probably the shortest possible transaction
  - but opens you up to inconsistent reads



# App Server

---

# App Server

---

- Process concurrency occurs in an application server
- Explicit concurrency control - threads, synchronization, etc...
  - We almost never have to deal with this

# Solving

---

# Solving

---

- We can launch a new process for each session
  - impractical
  - too slow

# Process-per-request

---

# Process-per-request

---

- Processes / threads are pooled
  - Each can handle multiple requests
  - but only one at a time
- Good concurrency control, good isolation

# Thread-per-request

---

# Thread-per-request

---

- Multiple threads in the same process
- Less resource intensive
  - more requests with the same hardware
- No isolation



# Choosing

---

# Choosing

---

- process-per-request is probably the easiest
  - good for a less experienced team
- I prefer thread-per-request
  - just requires a little more thought up front

# Session State

---



# Statelessness

---

# Statelessness

---

- Statelessness doesn't refer to an object without state
  - What it means is
    - an object does not retain state between requests

# Stateless

---

# Stateless

---

- The same Java object can be reused over and over again
- The same service method can be used over and over again

# Stateful

---



# Stateful

---

- Objects that maintain their state between user requests
- Can be a drain on memory resources

but...

---

# but...

---

- we can't avoid state
- because client interactions are “inherently stateful”
  - A users shopping cart is state, and state that they certainly want you to remember

# Session State

---

# Session State

---

- Data that is relevant only in the context of a particular user's session is considered session state
  - Lives only for the duration of the users session (or slightly beyond)

# Consistency

---

# Consistency

---

- Data in the users session might not be legal
  - They have removed their zip code from their address, but we haven't persisted (or validated) yet
- This is common in a users session

# Storing

---





# Storing

---

- There are many ways to store a users session
  - Client Session State
  - Server Session State
  - Database Session State



# Client Session

---

# Client Session

---

- Stores all session state with the client
  - URL encoded, cookies, serialized data
  - Just storing in memory (rich client)

# Server Session

---

# Server Session

---

- The server holds the data between requests
  - in memory
  - in a file
  - in a database table

# Database Session

---

# Database Session

---

- Slightly different from server session state stored in the database
- Instead of serializing the user's session you use (what I'm calling) shadow tables
  - These tables are like the master tables but only temporarily store data



# Client Session

---



# Client Session

---

- Totally impractical for the web when you have any significant amount of session information
  - you can make it work for a few fields
- This is because you have to transfer all of the information back and forth on every request
- Potential security problem - maybe you don't want to let the client see session data

# Isolation

---

# Isolation

---

- Maintaining isolation with database session state is difficult (according to our book)
- I think it is actually relatively easy to do - but with performance consequences

# Server State

---

# Server State

---

- So server state is probably our best option
- This is great if we have one application server
  - what about multiple servers?

# Session Migration

---

# Session Migration

---

- Allows a session to flow from one application server to another
  - works well when your session data is stored in the database
  - Can be done with in memory/file storage - but this is additional bandwidth between servers

# Server Affinity

---



# Server Affinity

---

- Force a user to use the same application server for their entire session
- This works but
  - bad for load balancing
  - horrible for redundancy (server crash and the users session is toast)

# What I think

---

# What I think

---

- User Server Session State with files (1 server) or database (1 or more servers) backing
  - Database here can be interpreted many ways:
    - Actual database
    - memcached
    - Not keeping session state, but re-reading current user state on each operation

# Real world

---

# Real world

---

- amazon.com uses some form of Database Session State
- Items are remembered between sessions

# Timeout

---

# Timeout

---

- What happens when the users closes their browser and doesn't shop anymore
- This can lead to memory bloat if we just wait for them to come back
  - so, sessions typically expire after some period of inactivity

# Client State

---



# Client State

---

- In the web - this has a place as well
- We typically shuffle a token back and forth (128, 160, or more unique bits)
- This is refereed to as a session ID



# Distribution Strategies

# Distributed Objects

---



# Distributed Objects

---

- In general
  - These are bad
  - Overused
  - Provide poor performance
  - Are not scalable architectures

# Remote / Local

---



# Remote / Local

---

- computers are the reason distributed objects don't work so well
  - A method call within a process is very fast
  - A method call between processes is markedly slower
  - A method call between processes on different machines, even slower



and...

---

# and...

---

- As a result how we deal with a remote object versus a local object is different
  - local: fine-grained interface. Individual getter methods for each field
  - remote: this model breaks down - I can't retrieve individual fields, but rather need to get as much as I can at once



SO...

---

# SO...

---

- the programming model fundamentally changes

# Tools

---

# Tools

---

- The tools provide transparency so it is not known to the caller that an object is local or remote
- but...
  - the programming model should be different, so we have to know

# Classes

---

# Classes

---

- A distribution strategy based around classes (as we know them) doesn't work

# First Law of Distributed Objects Design

---



# First Law of Distributed Objects Design

---

- “First law of Distributed Object Design: Don’t distribute your objects”





# Multiple CPUs

---

# Multiple CPUs

---

- Of course we want to use multiple processes, across multiple machines in order to scale
- Clustering is usually the answer
  - Run the full application on many machines
  - Everyone makes local calls, goes faster
  - Scalability is still present

# When you have to

---

# When you have to

---

- There are several places where you must distribute your processes
- The goal is to then minimize the distribute boundaries

# Client / Server

---

# Client / Server

---

- In traditional client / sever programming there is a clear distribution requirement

# Server / Database

---

# Server / Database

---

- The server and the database typically run on different machines
- You can go to extremes of running the database in the same process as the application (I've tried this)
- but... database servers are designed to be remote from the application and are usually optimized to that effect





# Web / Application

---

# Web / Application

---

- Sometimes these two jobs are split
- Don't split them if you don't have to

# 3rd Party

---

# 3rd Party

---

- Some 3rd party code can be compiled into your process, some cannot
- Usually can't do anything about this
  - hopefully they are coarse grained calls

# Just have to

---

# Just have to

---

- Maybe your higher ups have decided to force a distributed architecture
- Make your interfaces coarse grained

# Distribution Boundary

---

# Distribution Boundary

---

- We would still like to code with fine-grained objects
- The key is to use a Remote Facade to handle the coarse grained behavior
  - this is sort of a gateway to the remote process
  - I've found it useful even in local circumstances



# Interfaces

---

# Interfaces

---

- XML being sent over HTTP is the most common distribution technique at the moment
- Good interoperability

# Summary

---

# Summary

---

- We are going to build a framework and application side by side
- The best frameworks are extracted from real world applications