

# Offline Concurrency Patterns (Chapter 16)

---

Mike Helmick  
Large Scale Software Engineering  
Spring 2014

# Chapter 16

---



# Chapter 16

---

- Optimistic Offline Lock
- Pessimistic Offline Lock
- Coarse-Grained Lock
- Implicit Lock



# Optimistic Offline Lock



# Optimistic Offline Lock

---



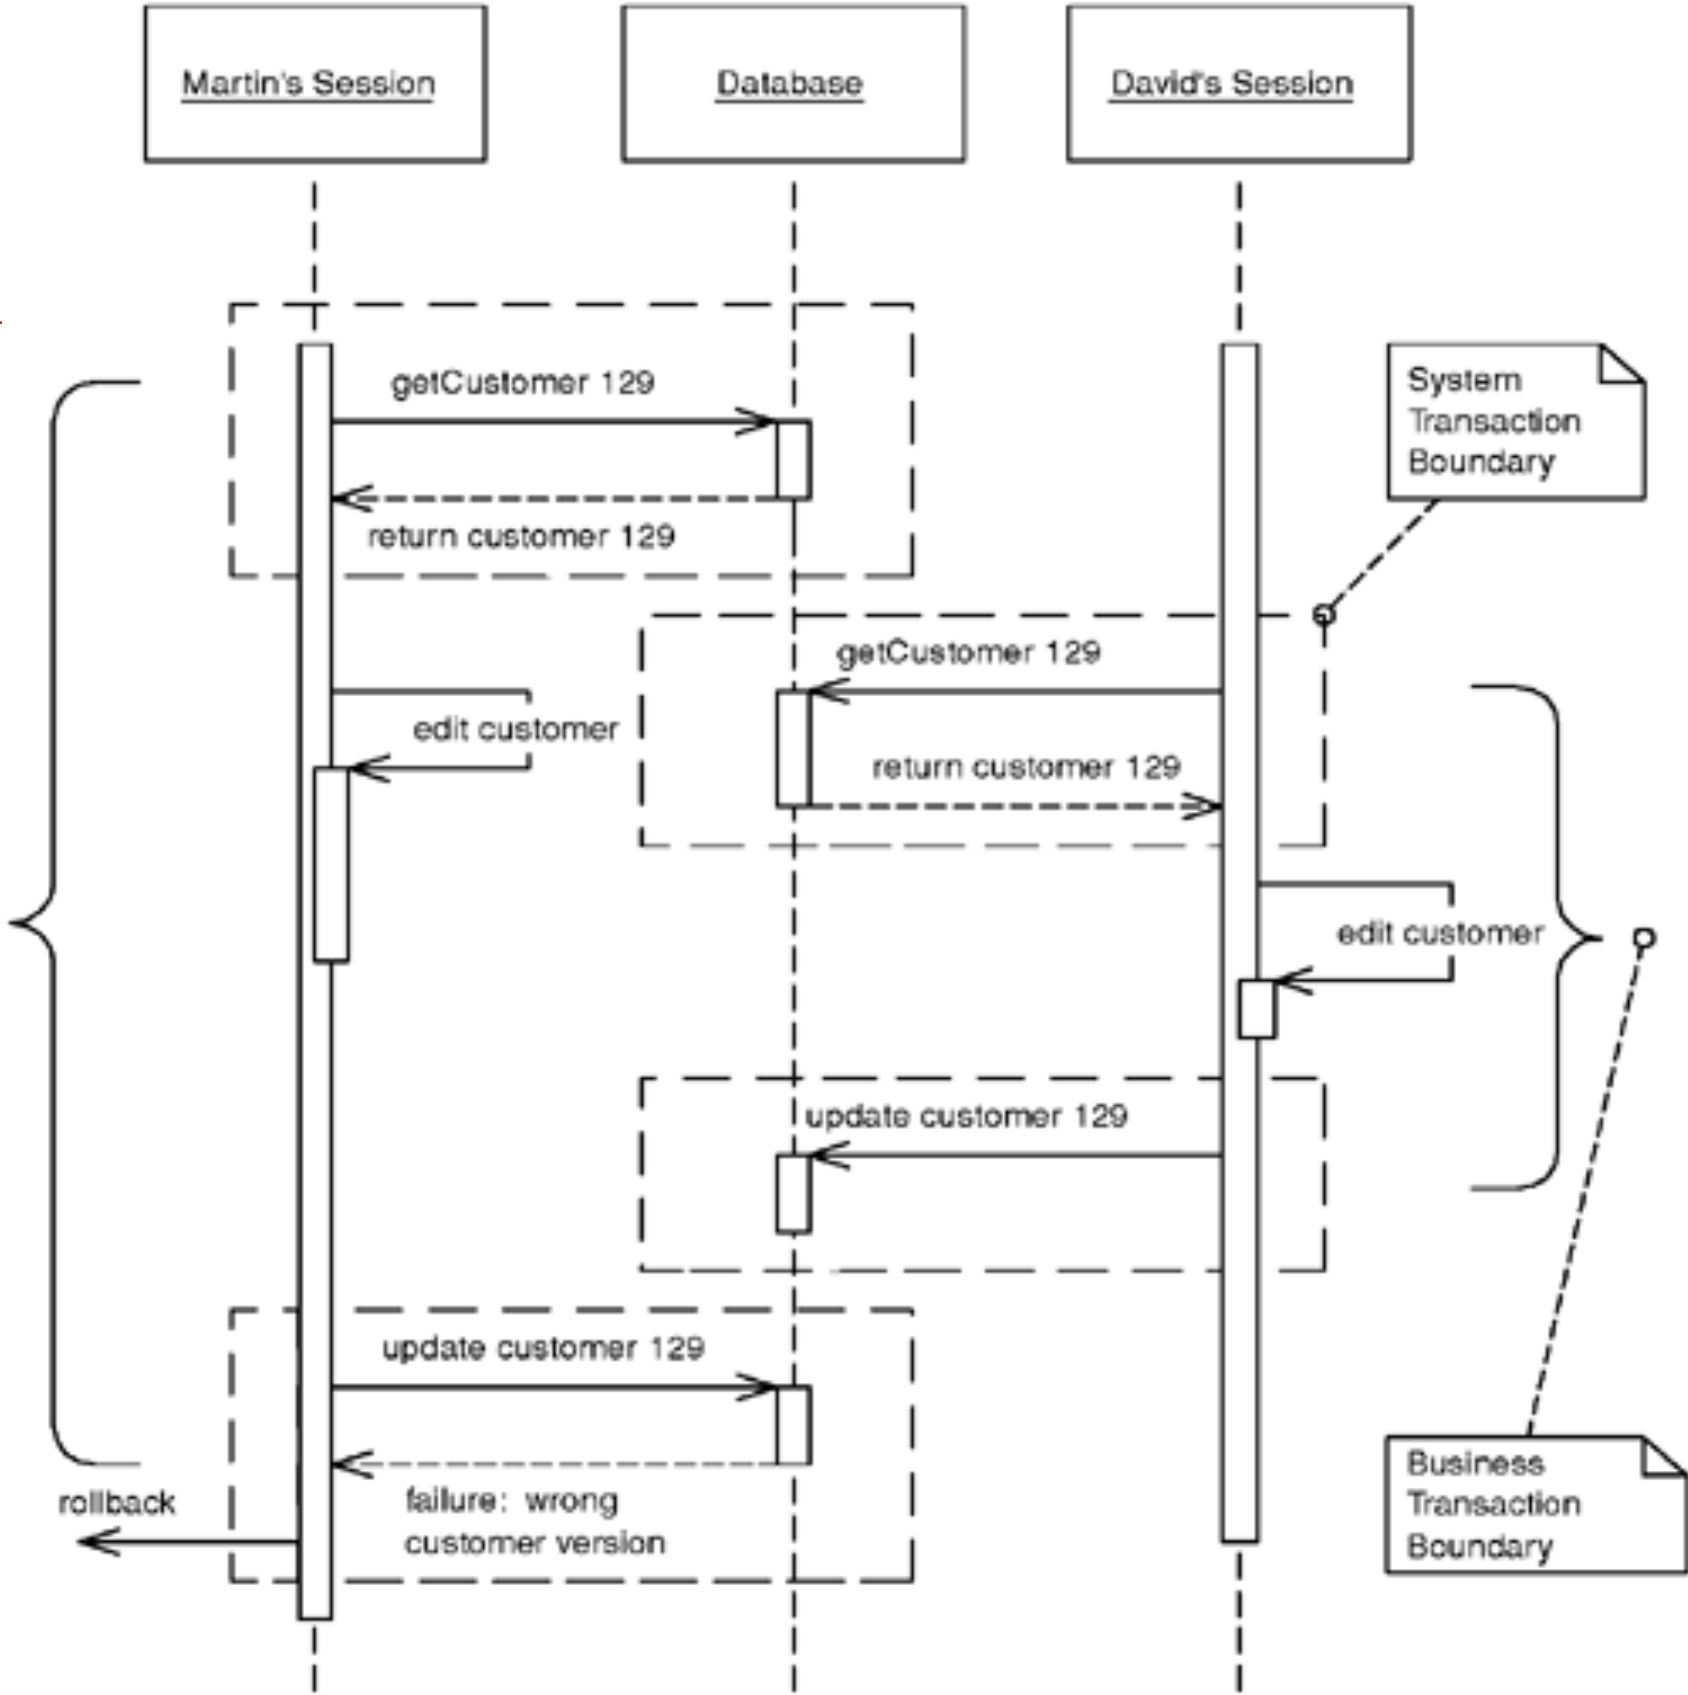
# Optimistic Offline Lock

---

- “Prevents conflicts between concurrent business transactions by detecting a conflict and rolling back transactions.”



# Optimistic Offline Lock



# Optimistic Offline Lock

---



# Optimistic Offline Lock

---

- Business transactions will normally span several system items
  - Some service calls will trigger several changes in the domain model
  - Two (or more) long running transactions occurring simultaneously

# Optimistic Offline Lock

---

# Optimistic Offline Lock

---

- Solves the concurrency problem by validating that the changes of this session do not conflict with other sessions
  - just before commit
- Assumes that the risk of conflict is low

# How It Works

---

# How It Works

---

- Optimistic Locks are not really locks at all
  - Validate that the data that was read is still the data that is in the database before the update
- Will ensure that no inconsistent data is written
  - but only after the transaction is finished processing

# How It Works

---

# How It Works

---

- Generally
  - this is implemented with a version number field on each row
  - compare the version number that was read on update
  - update the version number to a new one

# SQL

---



# SQL

---

- update table set ....., version\_number=?  
where version\_number=?
- using verNbr + 1 and verNbr as the bind variables
- Do this on updates and deletes

# Row Count

---

# Row Count

---

- In order to verify that the execution was successful, verify the row count
- For an update or delete of an identity row
  - the row count should be 1
- execute update returns a row count

# Audits

---

# Audits

---

- It can also be useful to store the user id of the person logged in when the row was last updated
- Then - you can also give the user a good error message
  - “Cannot update row - recently updated by X”

# Clock

---

# Clock

---

- Generally timestamp doesn't work
- why?

# Alternative

---



# Alternative

---

- Always use all fields in the where clause
  - then you don't need to have a version number column
  - this can cause major performance problems
  - and unnecessarily long where clauses

# Reads

---

# Reads

---

- Offline lock allows for inconsistent reads to occur
  - these are detected on updates
- but not detected on fields that are read and then used in calculations
  - so - a mechanism should be built to verify that the version number is still correct on commit

# Business

---

# Business

---

- Often concurrency is a business issue
- If a customer changes their service plan for instance
  - When their next billing cycle occurs which one do we use
    - the old one or the new one
    - do we have to use both



# SCM

---

# SCM

---

- Source code management systems (mostly) use optimistic offline lock
- generally OK because commit conflicts can be automatically (or manually) merged

# When to Use It

---



# When to Use It

---

- Most useful when the change of conflict in two database transactions is relatively low
- You don't want a user to continuously get version conflict errors
  - In my experience, I haven't seen this error occur twice in a row

# When to Use It

---

# When to Use It

---

- Easier to implement
  - doesn't have same runtime consequences as pessimistic offline lock
- You should always consider using optimistic lock
  - and escalate to pessimistic lock in certain circumstances

Example

# Pessimistic Offline Lock

# Pessimistic Offline Lock

---



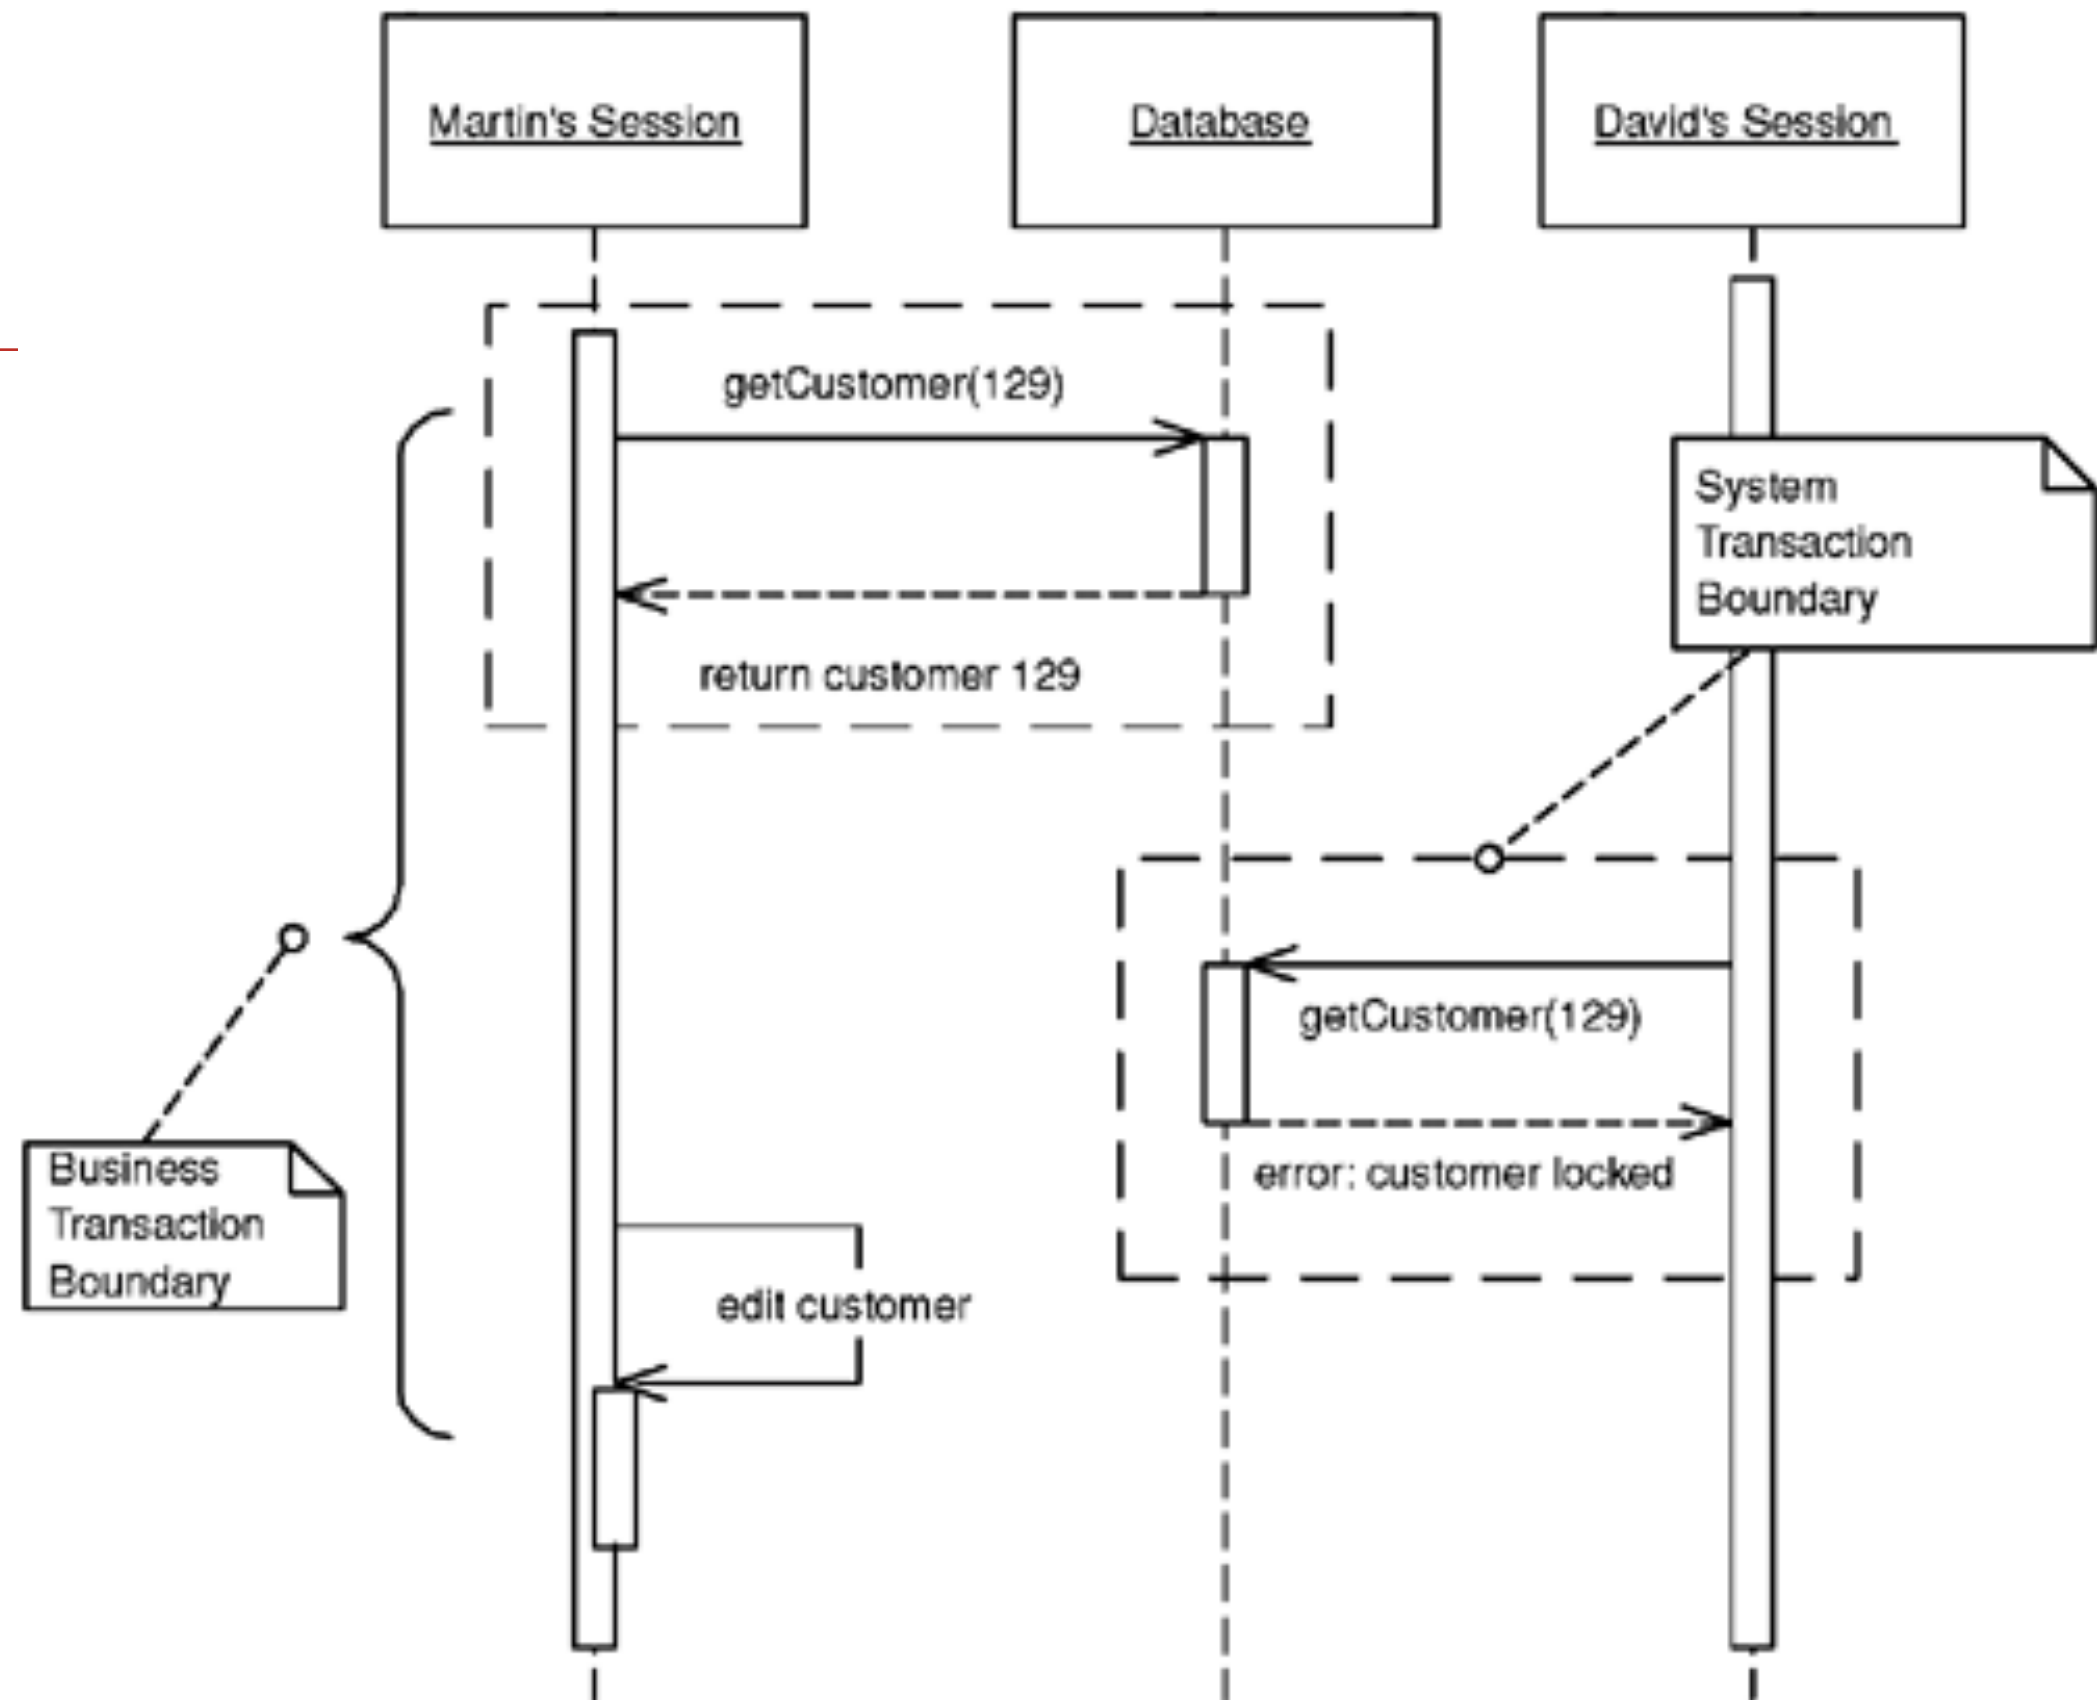
# Pessimistic Offline Lock

---

- “Prevents conflicts between concurrent business transactions by allowing only one business transaction at a time to access data.”



# Pessimistic Offline Lock





# Pessimistic Offline Lock

---

# Pessimistic Offline Lock

---

- Used for long business transactions that span multiple system transactions
- Optimistic offline lock allows a transaction to run to completion and fail at the last second
- Pessimistic lock prevents conflicts from occurring in the first place



# How It Works

---

# How It Works

---

- Implemented in 3 phases
  - determine types of locks
  - build the lock manager
  - define lock usage semantics

# Exclusive Write Lock

---

# Exclusive Write Lock

---

- Requires a business transaction to lock data only when editing
- Allows other processes to read locked data (and display it to the user)
  - but that user will get an error if they try to open the data for editing

# Exclusive Read Lock

---

# Exclusive Read Lock

---

- Used when you must always display the most accurate information
  - i.e. can't display data if another process has requested it for writing - or reading and still has the page open
- Severely restricts the concurrency of the system



# read/write lock

---



# read/write lock

---

- mutually exclusive: a single process can get a read lock or a write lock - no other process can get a lock if another has a write lock
- concurrent read locks are acceptable. As long as a read lock is in effect, anyone can open a read lock. Write locks block
- known as the readers / writers problem



# concurrency

---

# concurrency

---

- read/write locks improves concurrency
- major downsides
  - this is very difficult to implement
  - challenging in deciding which objects need read/write lock
  - challenging in creating an ordering that avoids deadlock



# To Consider

---



# To Consider

---

- Maximize concurrency
- Meet business rules
- Minimize complexity



# Pessimistic Lock

---

# Pessimistic Lock

---

- Can cause problems if implemented incorrectly
- Degrade system performance
- To many locks - escalating locks, etc....
  - in effect reduce you to a single user system



# Lock Manager

---

# Lock Manager

---

- The lock manager is the gateway for lock requests
- Must take in request and decide if it can be granted in the context of the business transaction, system transaction, and possibly the user (lock owner)



# Lock Manager

---

# Lock Manager

---

- A simple structure
  - a table that maps locks to owners
  - can be in memory or on the database
- Memory - use a synchronized singleton
- Database - works in clustered application environments

# Visibility

---

# Visibility

---

- Locks need to be private to the lock manager
- Business transactions should not be able to create their own locks

# Timing

---

# Timing

---

- In general
  - locks should be required before (as) the data is read
  - order of locks doesn't matter much for read only transactions
  - order is important for write locks!
    - deadlock



# What

---

# What

---

- You need to also decide what to lock
  - Best to lock the IDs - rather than the actual object
  - Then you can get the lock before reading the object (can't do this on search queries)

# Releasing

---

# Releasing

---

- Locks should be released when the business transaction completes
  - either success or failure
- Some locks can possibly be released early
  - a write lock obtained, but the the object doesn't change

# Errors

---

# Errors

---

- Since locks will be acquired early in the transaction
  - the user can be notified quicker
- Can possibly implement blocking locks
  - this will cause transactions to sit and wait until they are able to execute

# Database

---

# Database

---

- Storing your locks in the database causes some problems
  - The lock table itself needs to have concurrency control
  - Can help to use database locks in this instance



# Blocking

---

# Blocking

---

- Using something like “SELECT FOR UPDATE” can lead to deadlock
  - Where as offline lock will give an error in this situation

# Timeouts

---

# Timeouts

---

- the last thing to consider is lock timeout
- occasionally a client will crash after locks have been obtained
  - some reasonable timeout should exist

# When to Use It

---

# When to Use It

---


- When the chance of conflict is high
- When you want a client to never sit and wait
  - i.e. the get quicker error messages
- When you have long running transactions

Example

# Coarse-Grained Lock

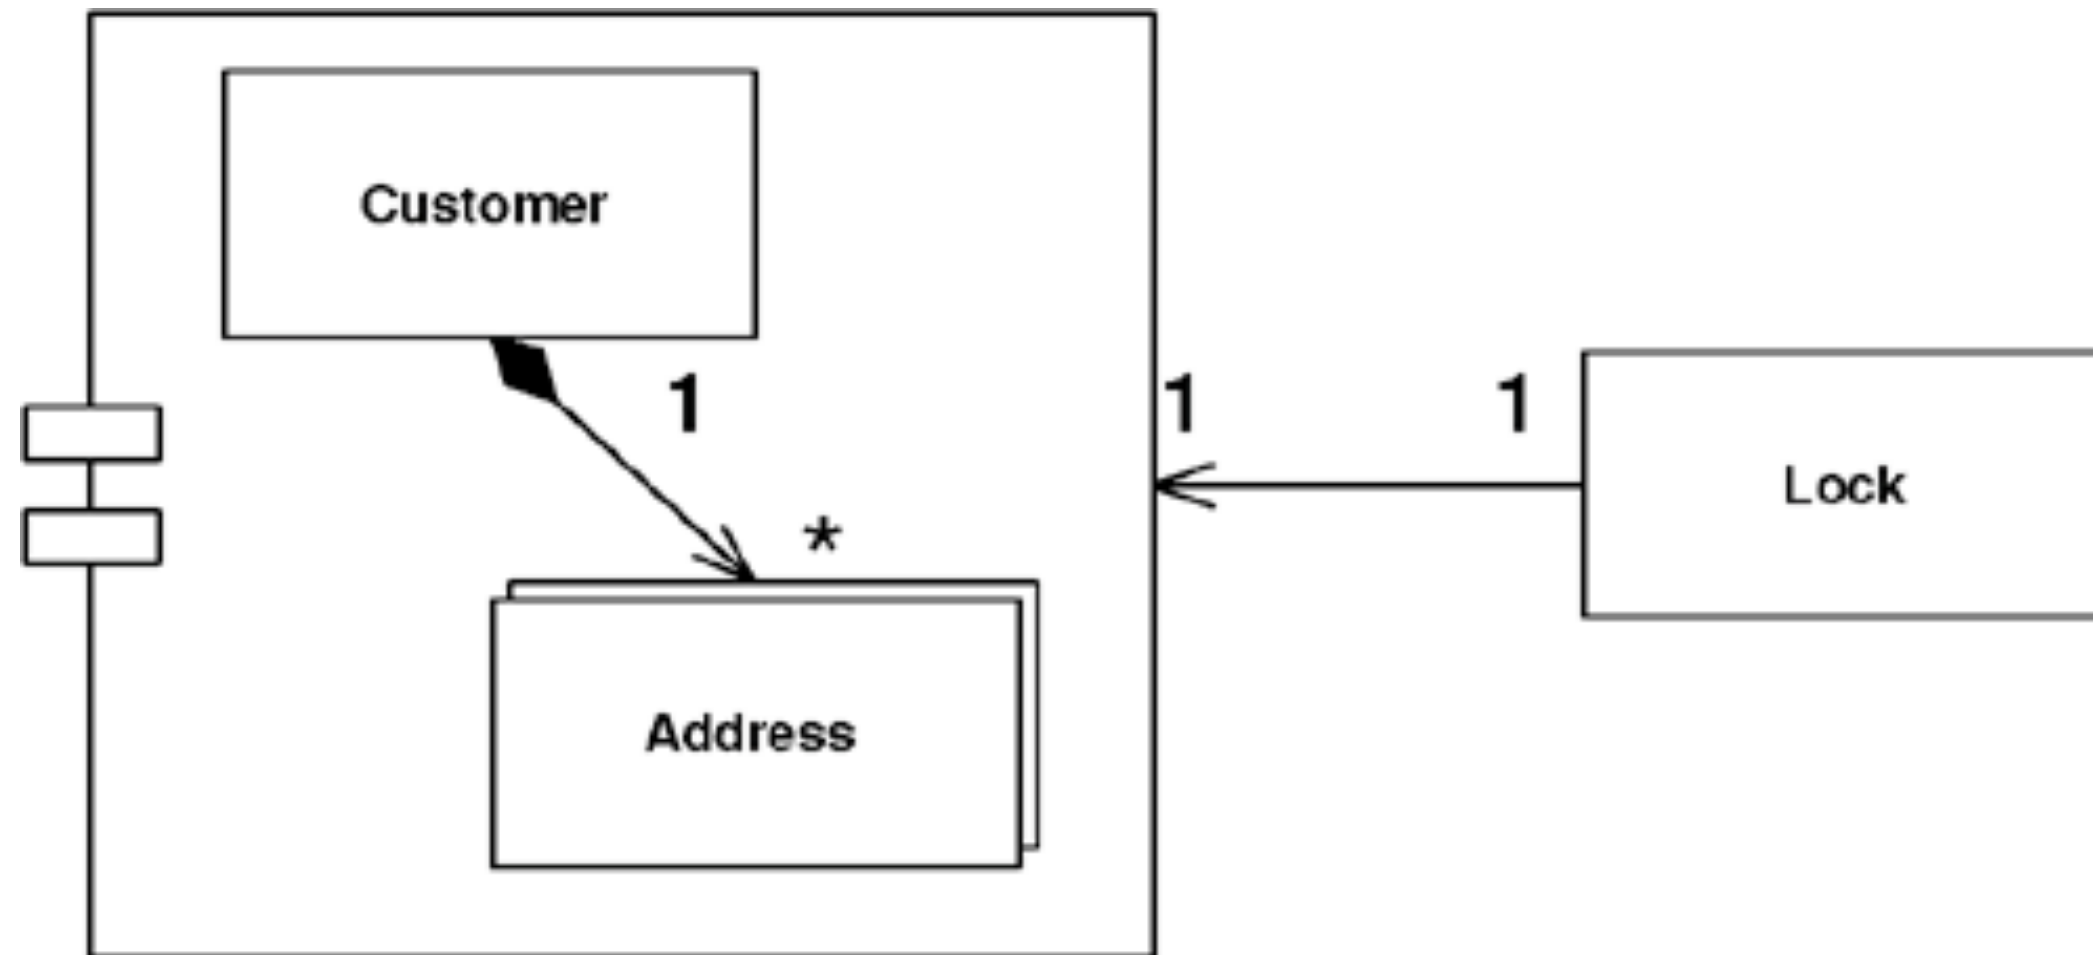


# Coarse-Grained Lock



# Coarse-Grained Lock

- “Locks a set of related objects with a single lock.”



# How It Works

---

# How It Works

---

- Create a single lock point for a group of related objects
  - Then only one lock is necessary, rather than several

# Optimistic

---

# Optimistic

---

- Use a shared version number for all records in the set
  - when any record in the group is used the shared lock is checked

# Pessimistic

---

# Pessimistic

---

- Every group member shares the same lock id
- Again - shared version number / object is a good candidate for a shared pessimistic lock id



# Aggregate

---

# Aggregate

---

- Data groups have been defined as an aggregate
  - has a root object
    - the lock boundary for the entire group
- Coarse grained lock mediated through the root is appropriate in this case

# Aggregate

---

# Aggregate

---

- Requires that subordinate objects to know the root of the hierarchy

# When to Use It

---

# When to Use It

---

- When business requirements call for it
- The act of locking is quicker - since the amount of records included in the lock is larger
  - so less locks are needed

# Chubby

---

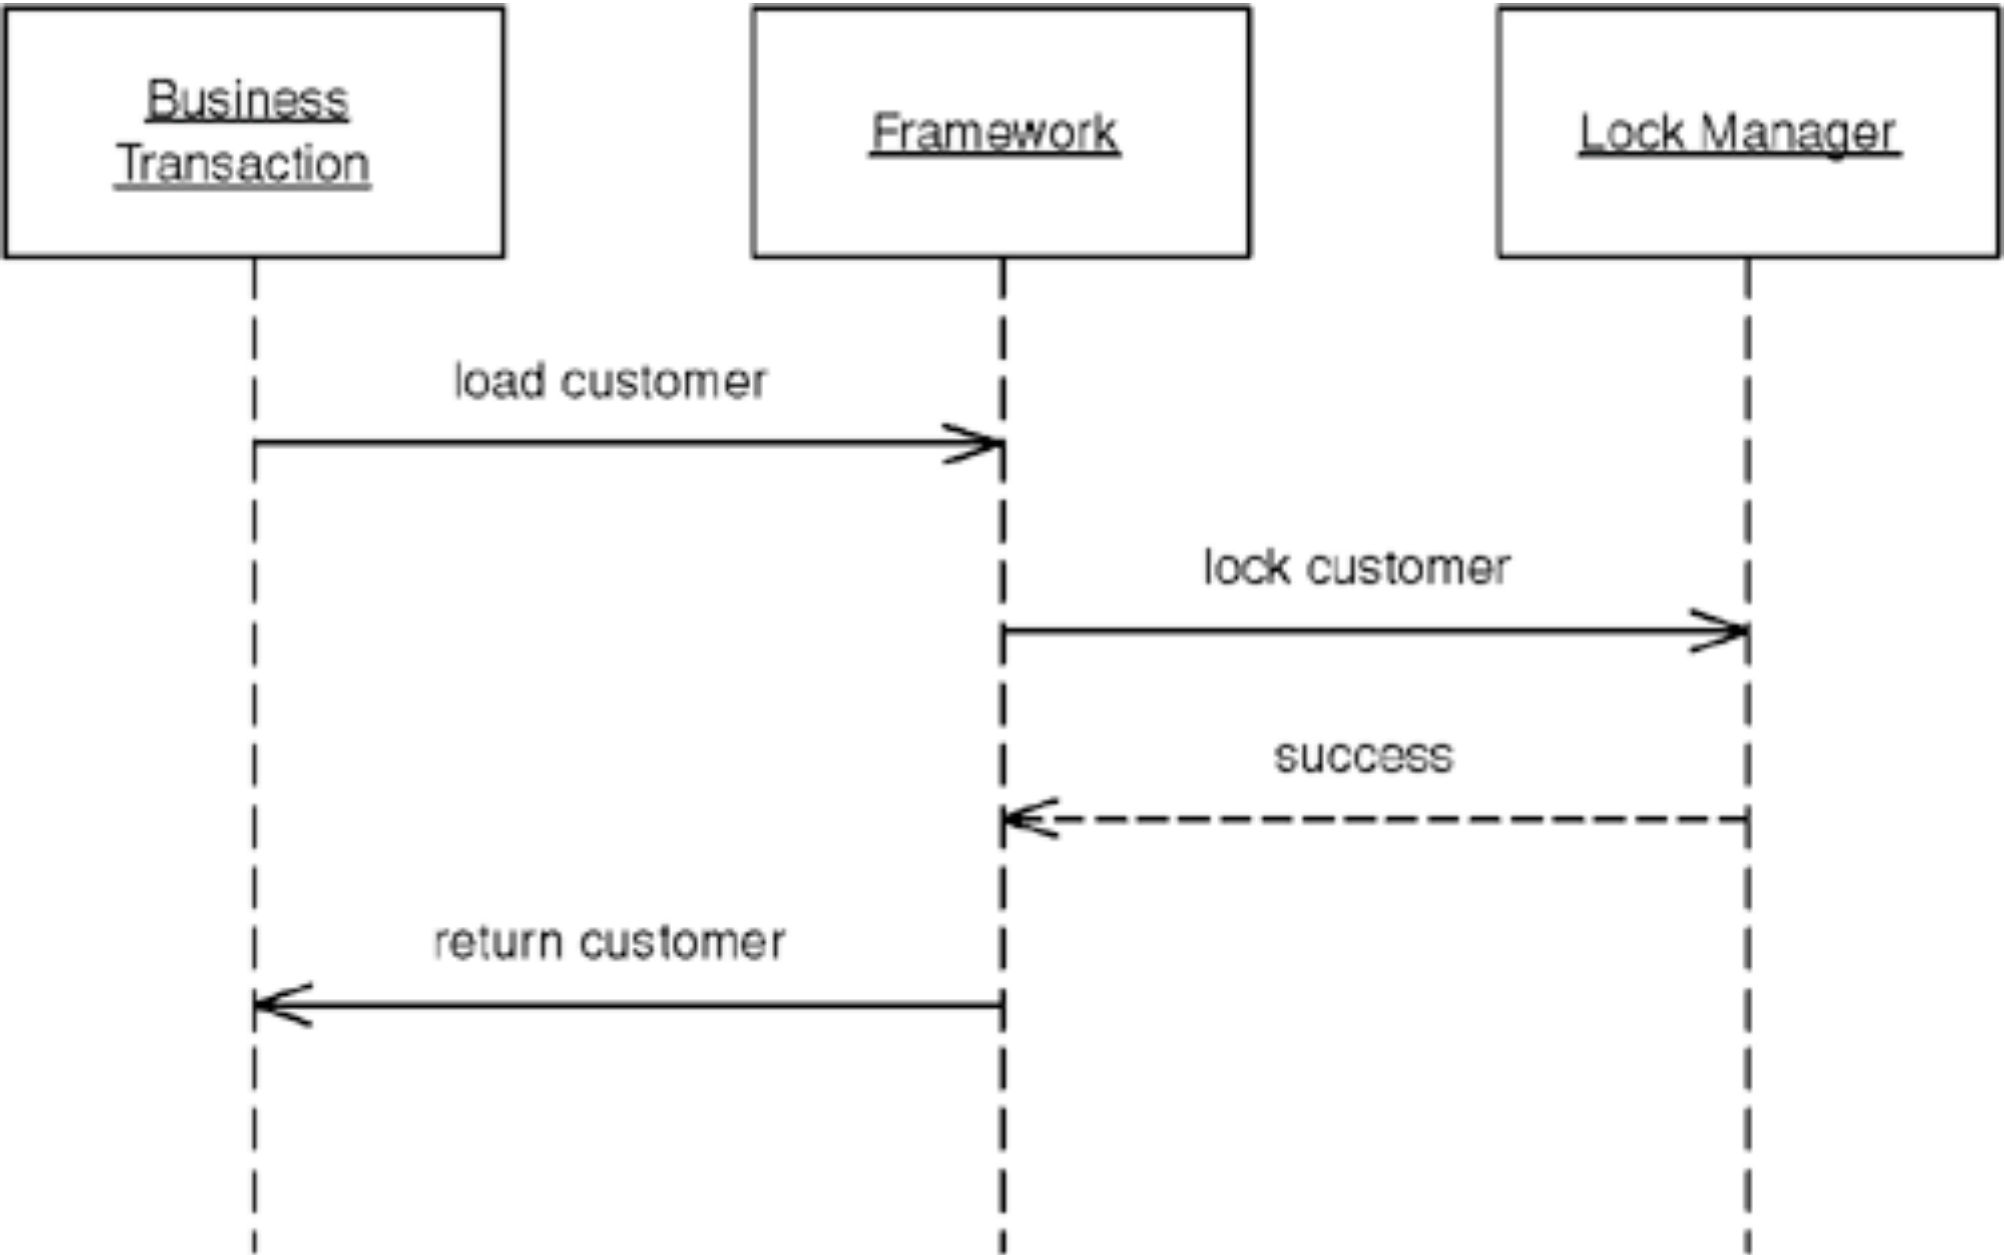
- This is where a distributed lock manage, like Chubby, will be useful

# Implicit Lock



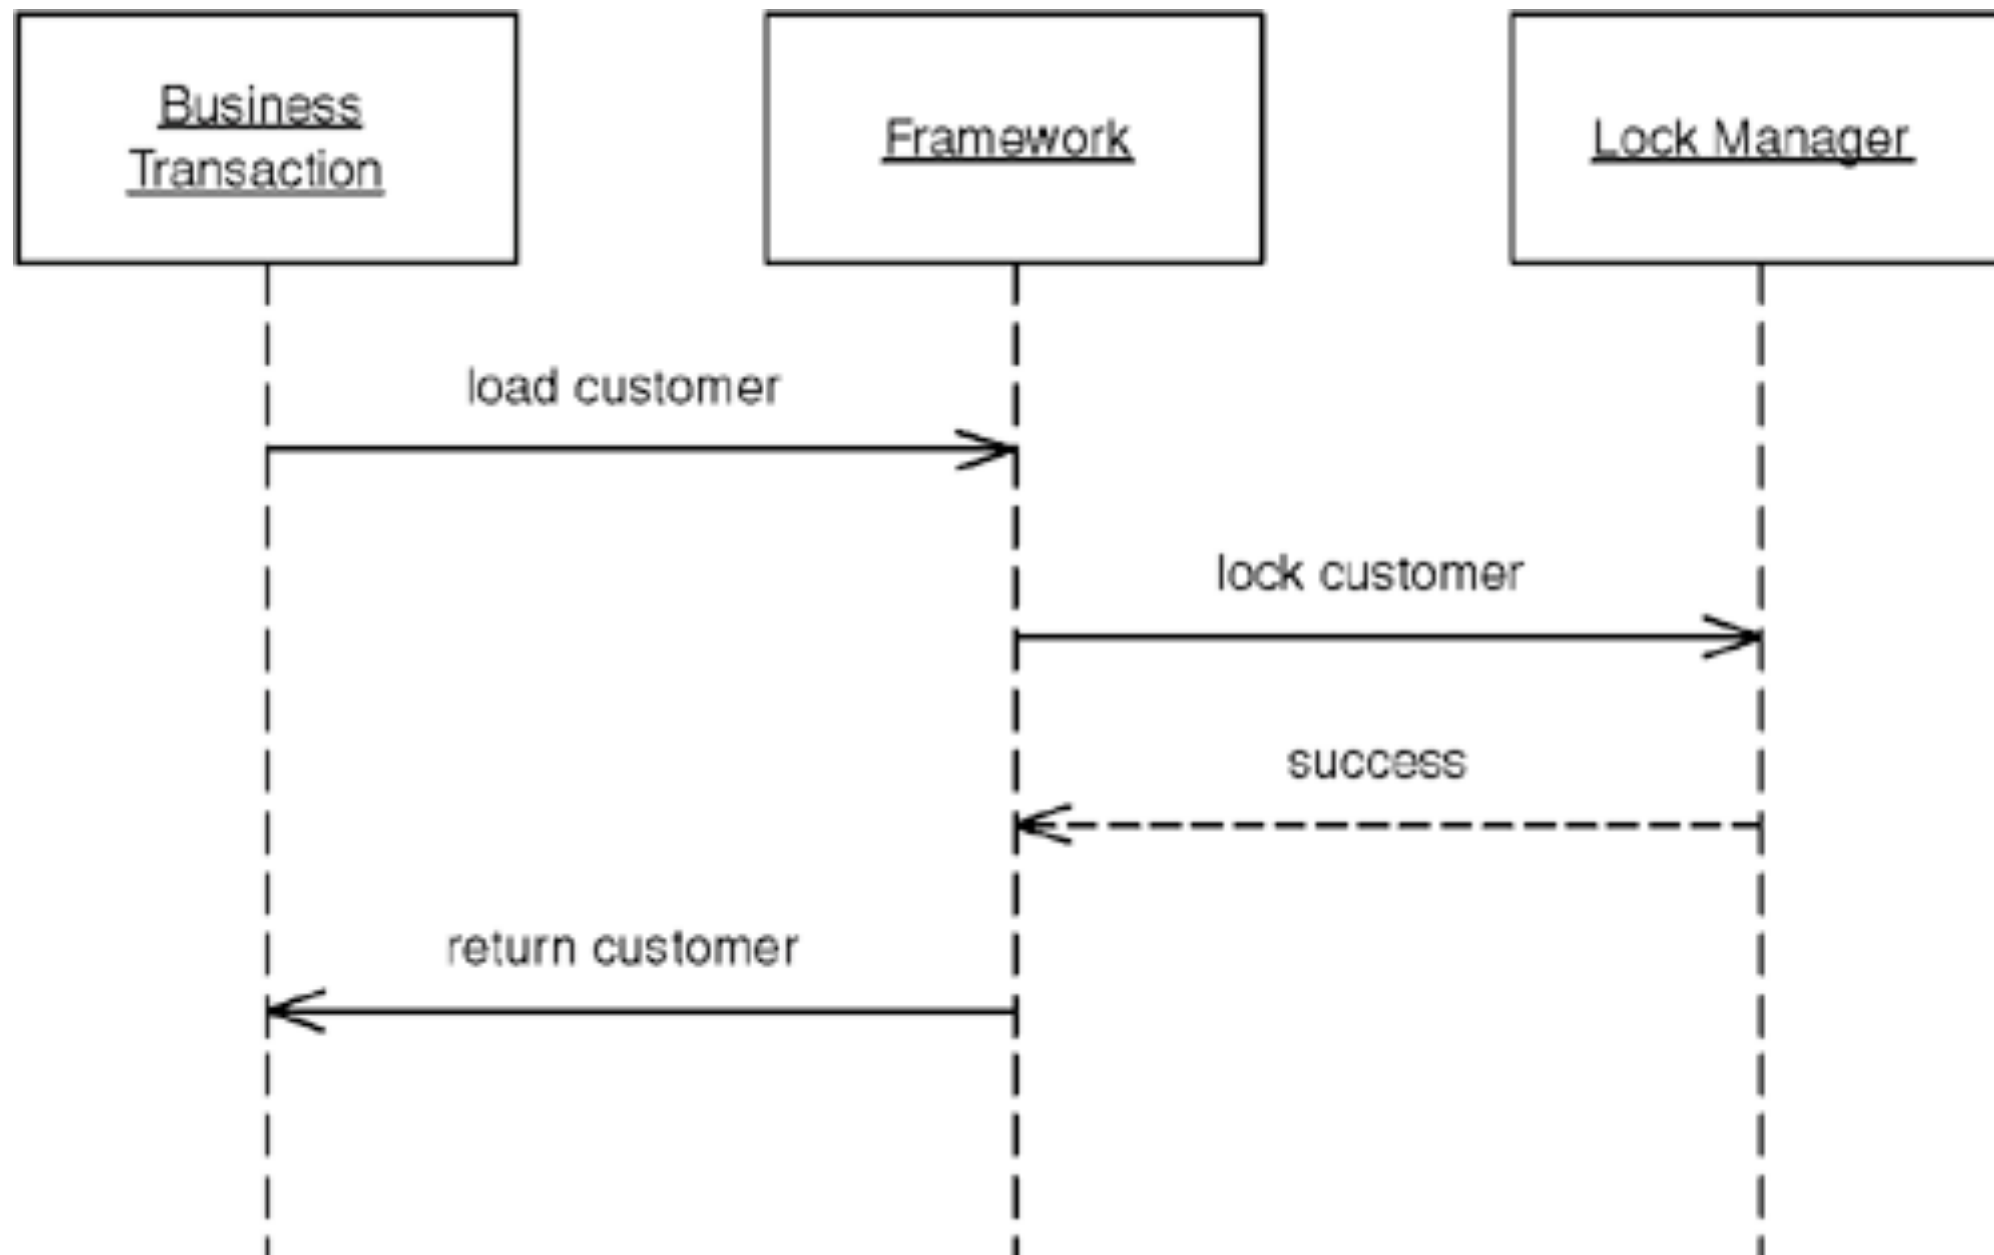


# Implicit Lock



# Implicit Lock

- “Allows framework or layer supertype code to acquire offline locks.”



# Implicit Lock

---

# Implicit Lock

---

- A successful locking mechanism is always in force
  - Best of locking can be done automatically
  - If developers must explicitly code locks, they are likely to forget



# How It Works

---

# How It Works

---

- Implicit lock
  - embedded in your code so that locks can not be circumvented in any way in the code
  - obviously direct SQL on the database can get around this

# How It Works

---

# How It Works

---

- Best if you can hook this into your data source layer framework



# Concurrency

---

# Concurrency

---

- Always using implicit lock can cause problems
  - It is possible that you want to move the locking out of the technical area into the domain area

# Danger

---

# Danger

---

- If using blocking pessimistic locks
  - the developers need to read everything in the correct order

# When to Use It

---

# When to Use It

---

- Whenever you use locks
- One missed lock cause the whole application to have the potential for failure